

# **Tips for Giving Clear Talks**

**Kayvon Fatahalian**

---

**Updated May 2021**

# My motivation

- I have found I give nearly the same feedback over and over to students making talks
  - It is not profound feedback, it is just application of a simple set of techniques and principles that are consistently useful when making talks
- I am hoping that:
  - These slides serve as a useful checklist that you can refer to vet your own talks before delivering them to others
  - I still make these mistakes all the time when creating first drafts of talks

# **Why give talks?**

# **Two reasons to give talks**

- 1. Convey what you have learned in your research to your peers**  
**(Don't tell the room what you did, tell the room the most important things they should know, but probably do not.)**
  
- 2. Get feedback from others to advance your own research goals**  
**(Goal: put smart people in the best position to help you.)**

# **Two BAD reasons to give talks**

## **1. You believe it is a reward for:**

- Finishing a project**
- Getting a paper accepted**
- Being clever enough to get a good result**

## **2. You were signed up to speak in lab meeting or in CS197 (Presenting is an opportunity, not an obligation.)**

# Consider the costs of a bad (unclear) talk

- To the audience:

- 1 hour x 50 people = 50 person-hours of work
- General unhappiness

- To you:

- Missed opportunity for feedback or quality discussion
- Missed opportunity to identify new collaborations
- Diminished impact of your work

# **Benefit TO YOU of a good (clear) talk**

- **Non-linear increase in the impact of your work**
  - **Others are more likely to remember and build upon your work**
  - **Others are more likely to adopt your ideas**
  - **Others are more likely to come up to you after the talk**
- **Clarity is highly prized in the world: the audience will remember clear communicators**
  - **“Hey, that was a great talk yesterday... are you looking for a job anytime soon?”**
  - **“Hey, that was a great talk, I’m working on something that you might find helpful.”**

# Roadmap

- The rest of this talk is structured as a list of principles and tips
  - It is not a comprehensive guide to making a talk
- So let's get started...

# Who painted this painting?



Salvador Dali (age 22)

**My point: learn the basic principles before you consciously choose to break them**



# **Tip 1**

**Identify your audience**

# Strive for clarity for that audience

- You should aim for your target audience to understand everything you say in a talk (if they won't understand it, why are you saying it?)
- This means you have to put yourself in your audience's shoes
  - Even if you are targeting experts (e.g., your research advisor or peers), experts haven't been thinking about your problem 60 hours a week
  - The ability to analyze your own talk from the perspective of others is a skill ~~young researchers~~ all researchers (including professors) struggle with tremendously.

Tip: recite a sentence out loud to yourself. Do you really expect someone who has not been working with you everyday on the project to understand what you just said? \*

\* I find hearing myself say something makes it easier to parse it from an audience's perspective.

# **Tip 2**

**A good principle for any talk (or paper):**  
**“Every sentence matters”**

**What are you trying to say?**

**What technical story are you trying to tell?**

**What is point you are trying to make?**

**Is what you just said making that point? (If not, remove it)**

**If you can't justify how it will help the listener, take it out.**  
**If it won't be understood, take it out.**

**An example of applying “every sentence matters”:  
The talk intro**

# Intros and background are often very poor

- Too many talks have introductions and related work that do not suggest much thought was put into them
  - Which is bad — those are critical parts of the talk!
  - Example of useless intro to a graphics audience: many people in computer graphics know that global illumination is important to realistic image synthesis. Don't have to say it.
- Unlike a paper, related work in a talk does not exist for academic completeness
  - No one cares if citations are comprehensive (will address this later)
- They intro/related work in a talk exist to set the context for the technical components of the talk. Specifically...

The goal of the intro and background is to tell the listener: “In this talk, here is the way I want you to think about the problem I am trying to solve.”

# Bad example 1

■ Never ever, ever, ever do this!

The slide has a blue header bar with the word "Outline" in white. Below the header is a white content area with a black border. The content area contains the following bulleted list:

- **Introduction**
- **Related Work**
- **Proposed System Architecture**
  - ❖ Basic design decision
  - ❖ Dedicated hardware for T&I
  - ❖ Reconfigurable processor for RGS
- **Results and Analysis**
- **Conclusion**

This slide just told me this talk will have an introduction and a conclusion. :-(

# Bad example 2

Who is the audience for this? (how does this benefit them?)

**Discrete Methods**

- Remove/Add Pixels
- Related work\*:
  - Seam Carving [SIGGRAPH 07]
  - Improved Seam Carving [SIGGRAPH 08]
  - Shift-Map Image Editing [ICCV 09]

\* non-exhaustive

From: Avidan et.al., Seam Carving for Content-Aware Image Resizing

From: Patch et.al., Shift-Map Image Editing

(a) (b)

(c) (d) (e)

## ■ Experts on the topic?

- They likely know these papers exist. These slides don't tell them what about these papers is most relevant to this talk

## ■ Non-experts?

- They won't learn the related work from these two slides

**Continuous methods**

- Find continuous transformation
- Warp/deformation grid
- Related work\*
  - Non-homogenous warping, ICCV 07
  - Streaming video, SIGGRAPH 09
  - Shrinkability Maps for Content-Aware Video Resizing, PG 09
  - Robust Image Retargeting via Axis-Aligned Deformation, EG 12



Generated with the Streaming Video approach [KRA09]



From: Robust Image Retargeting via Axis-Aligned Deformation

retargeting to 200% width using axis-aligned deformations

**The goal of the intro and background is to tell the listener:  
“During this talk, here is the way I want you to think about the  
problem I am trying to solve.”**

**An excellent strategy to catch the audience’s attention and frame  
the story is to make them aware that there is something they  
didn’t know they didn’t know.**

**(“You might think you know this, but here’s a new angle on it”)**

## Intro example #1: [Mullapudi 2019]

The talk intro simply asked the audience to consider the following question:

*Why are we working so hard to train general ML models that work in many different situations?*

*Why don't we just use simple, less general ML models and retrain them constantly for the specific conditions at hand?"*



Here's the sequence used to do that...

# This is what a traffic camera sees

It seems like it should be easy to train a simple, low-cost ML model to detect cars/pedestrians/etc in this one scene... right?



# Problem: distribution shift

But any one scene can change dramatically over time.

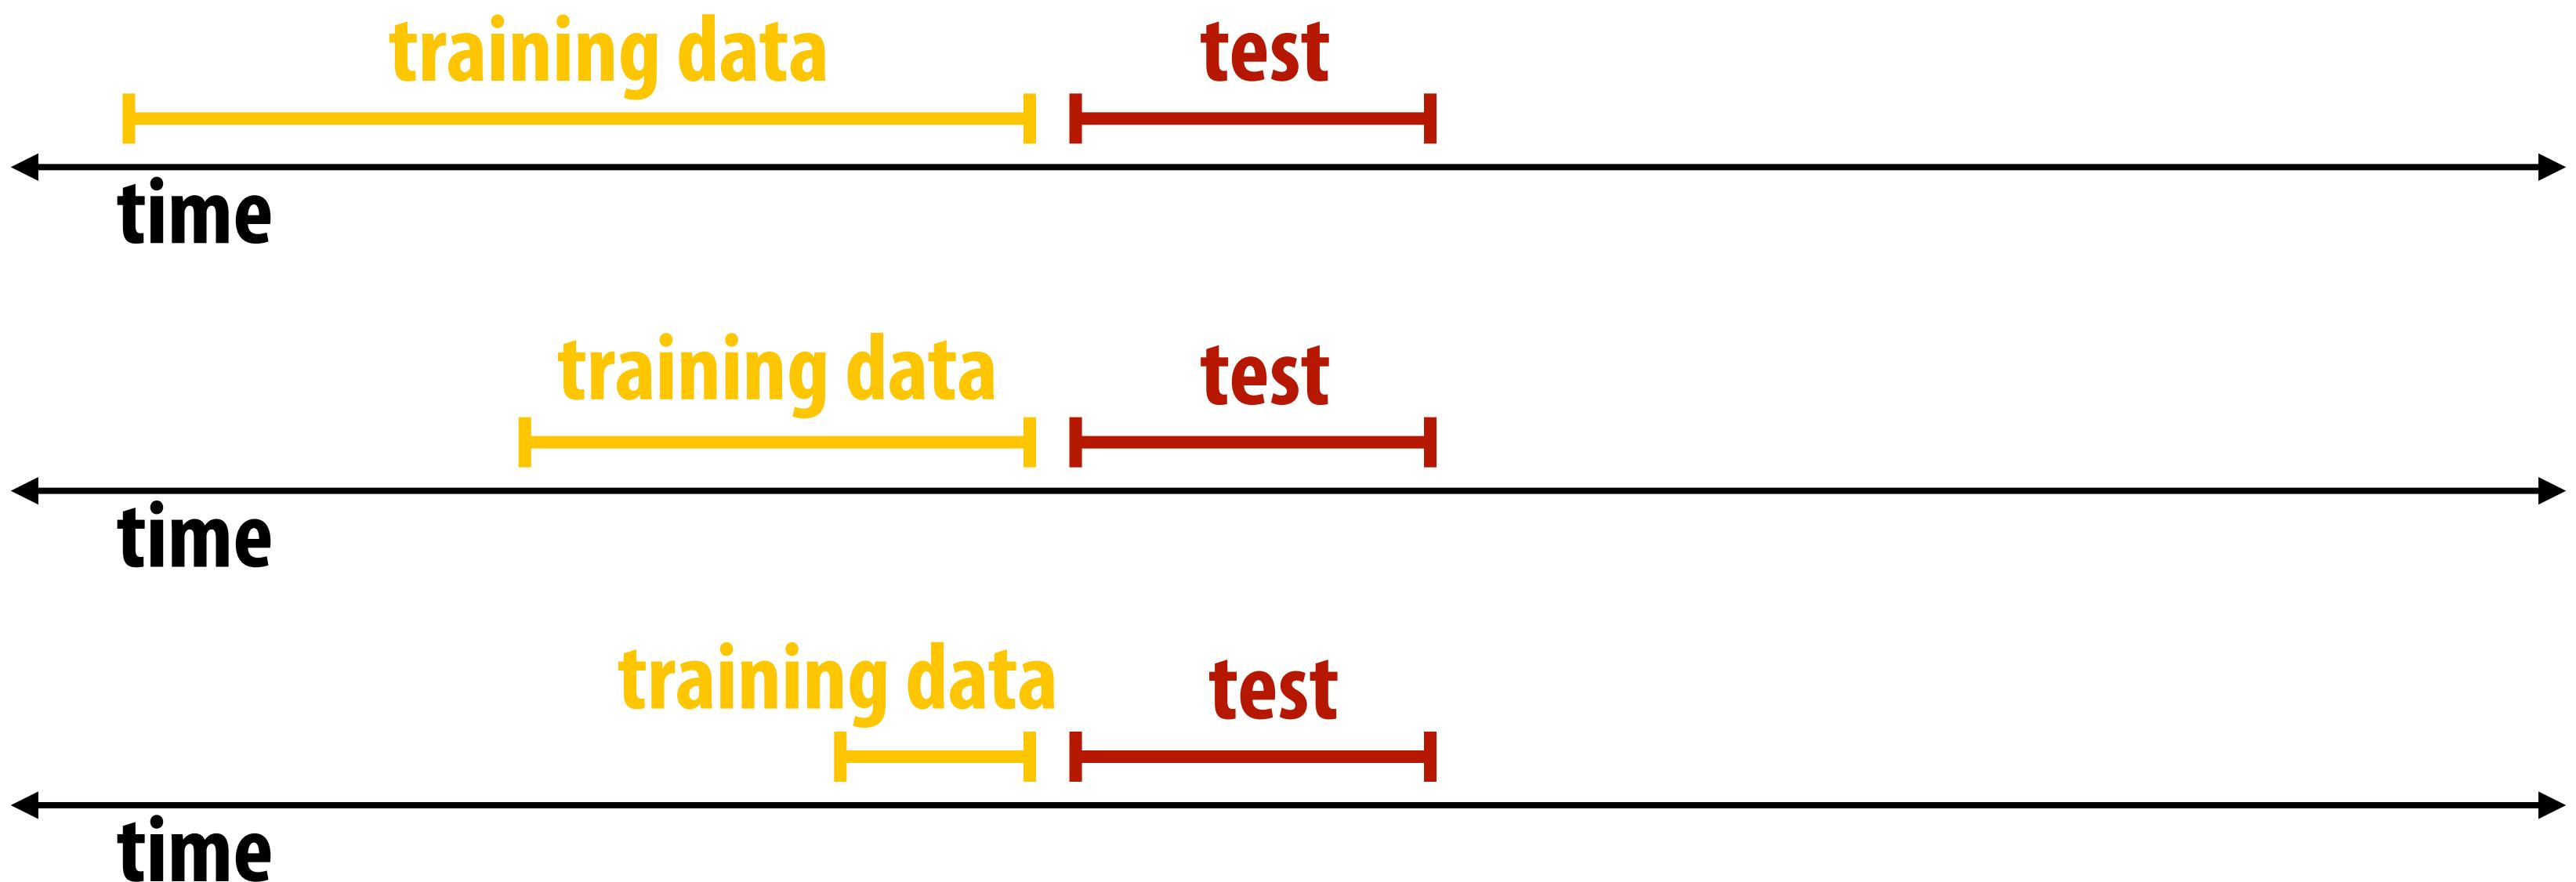
We need to acquire a training set for this diverse set of conditions... right?



Weather, time-of-day, types of vehicles in view, etc...

# Consider this experiment

- Plop camera down in a new environment
- We want a specialized model for processing the stream from this camera
- How much training data is needed to train an accurate model?



Answer: the small amount of recent training data worked best.

# Let's embrace continuous adaptation of ML models to the conditions at hand

- Our premise: low cost efficient models can retain high accuracy for complex tasks in challenging environments *if they are continuously specialized to the contents of video streams*
- A.k.a. Don't worry about carefully sampling everything you might see to create a good training set, just make sure you can adapt quickly online when you see it
- Audience member: “But how do I adapt an ML model quickly on the fly with little data or supervision? That sounds hard.”
  - Ah, that’s the point of the talk...
  - This person has something to say, maybe I shouldn’t check email.

**Now that we've framed the story, the related work can now be discussed in the context of this framing.**

**Establishing this framing is **the primary value** of the talk intro.**

**In this example: how have other folks tried to quickly adapt ML models to new contexts?**

*The "meta learning" field tries to train models that can be quickly adapted later using a few examples...*

*There's a whole body of work on "domain adaptation": taking a model that works well in context X and modifying it to work well in Y...*

*Instead we adopt an approach based on "model distillation", training a low cost student model to mimic the output of a general purpose teacher model.*

# **Tip 3**

**Set up the problem:  
establish inputs, outputs, and constraints  
(goals and assumptions)**

# Establish goals and assumptions early

- Given these inputs, we wish to generate these outputs
- We are working under the following constraints
  - Example: the outputs should have these properties
  - Example: the computer graphics algorithm...
    - Should run in real time
    - Should be parallelizable so it can run on a GPU
    - Should not require artist intervention to get good output
  - Example: the system...
    - Need not compile all of Python, only this subset that we care about...
    - Should realize about 90% of the performance of hand-tuned code, with much lower development time

# Why is knowing the goals and constraints important?

**Your contribution is typically a system or algorithm that meets the stated goals under the stated constraints.**

**Understanding whether a solution is “good” requires having this problem context.**

# **Tip 4**

**Show, don't tell**

**It's much easier to communicate with  
figures/images than text**

**(And it saves the speaker a lot of work explaining)**

# **Example:**

- In a recent project, we asked the question... given enough video of tennis matches of a professional athlete, could we come up with an algorithm for turning all this input video into a controllable video game character?

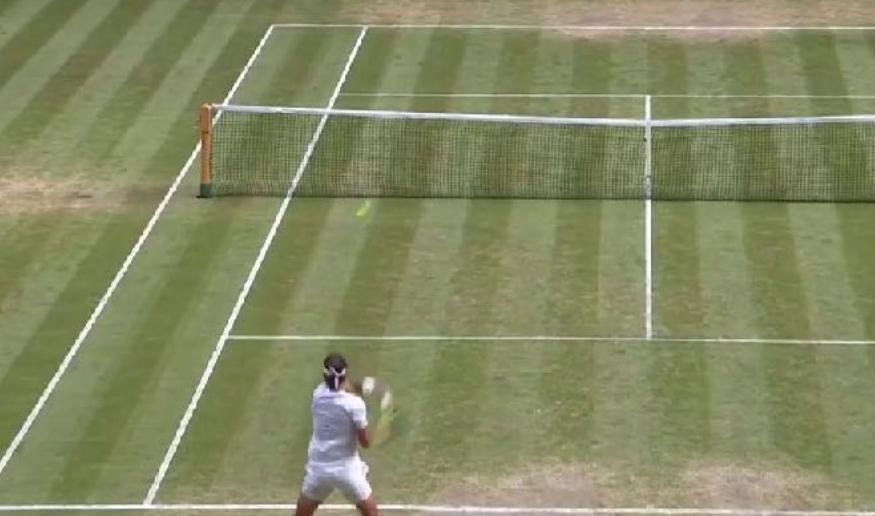
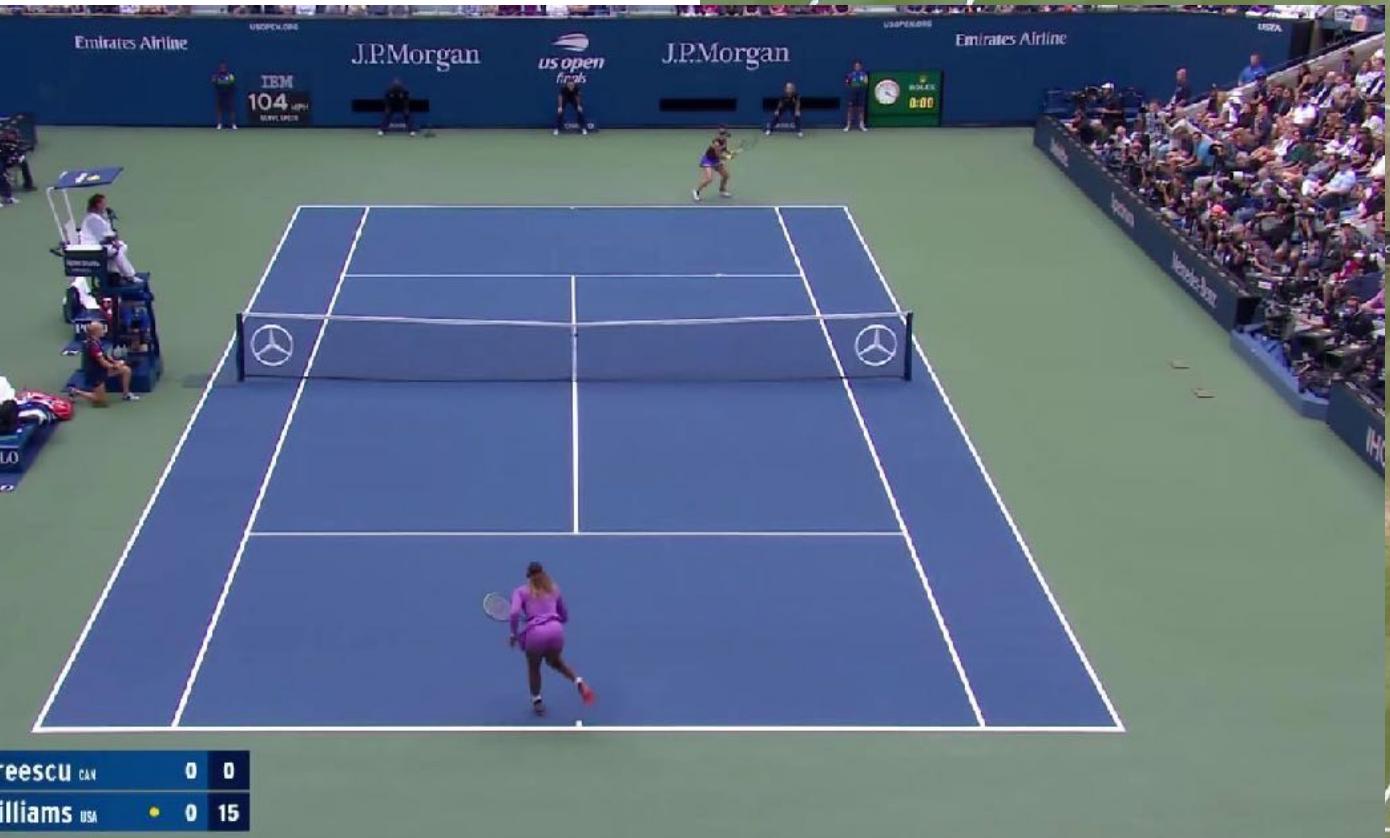
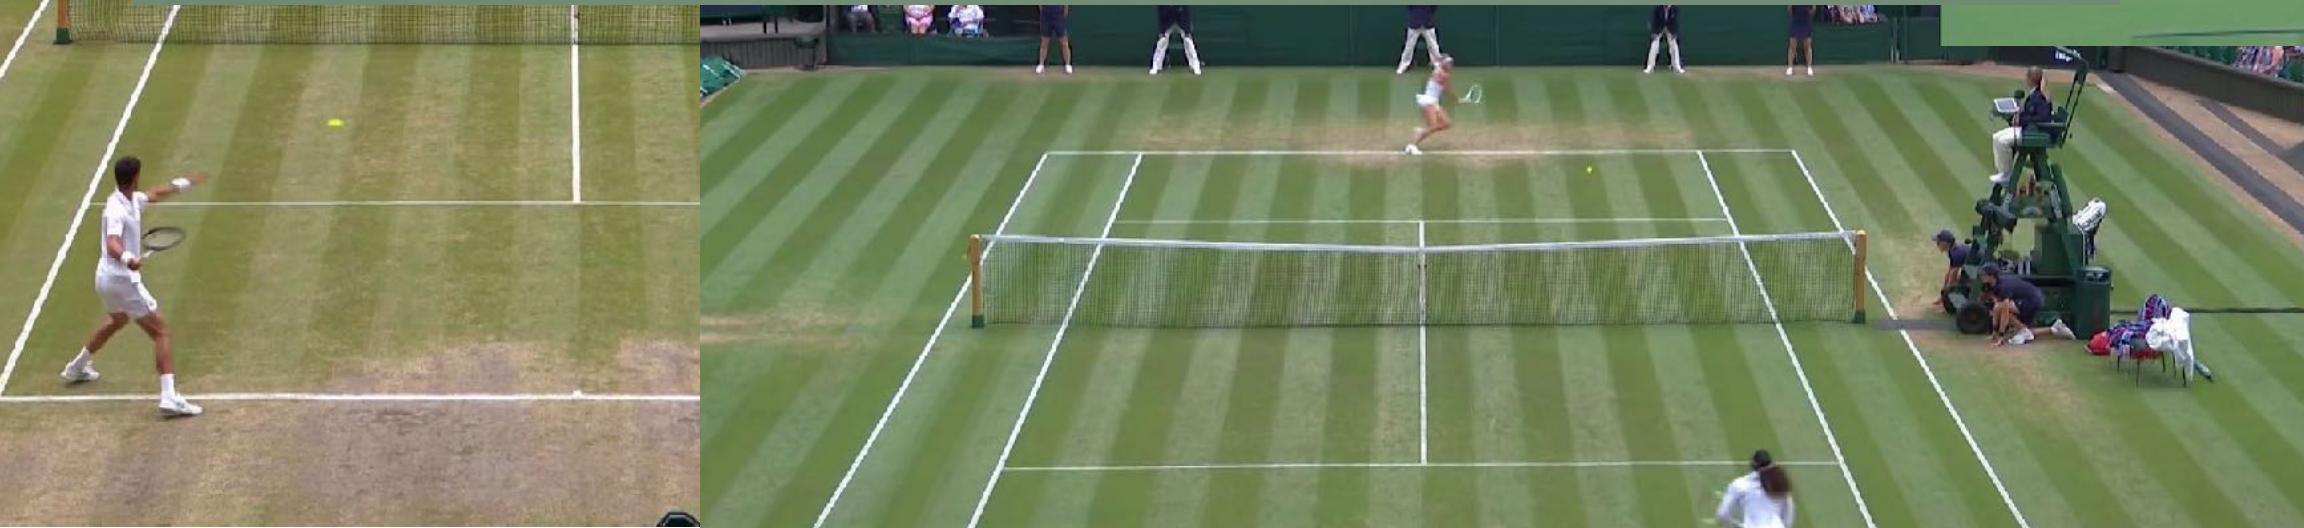
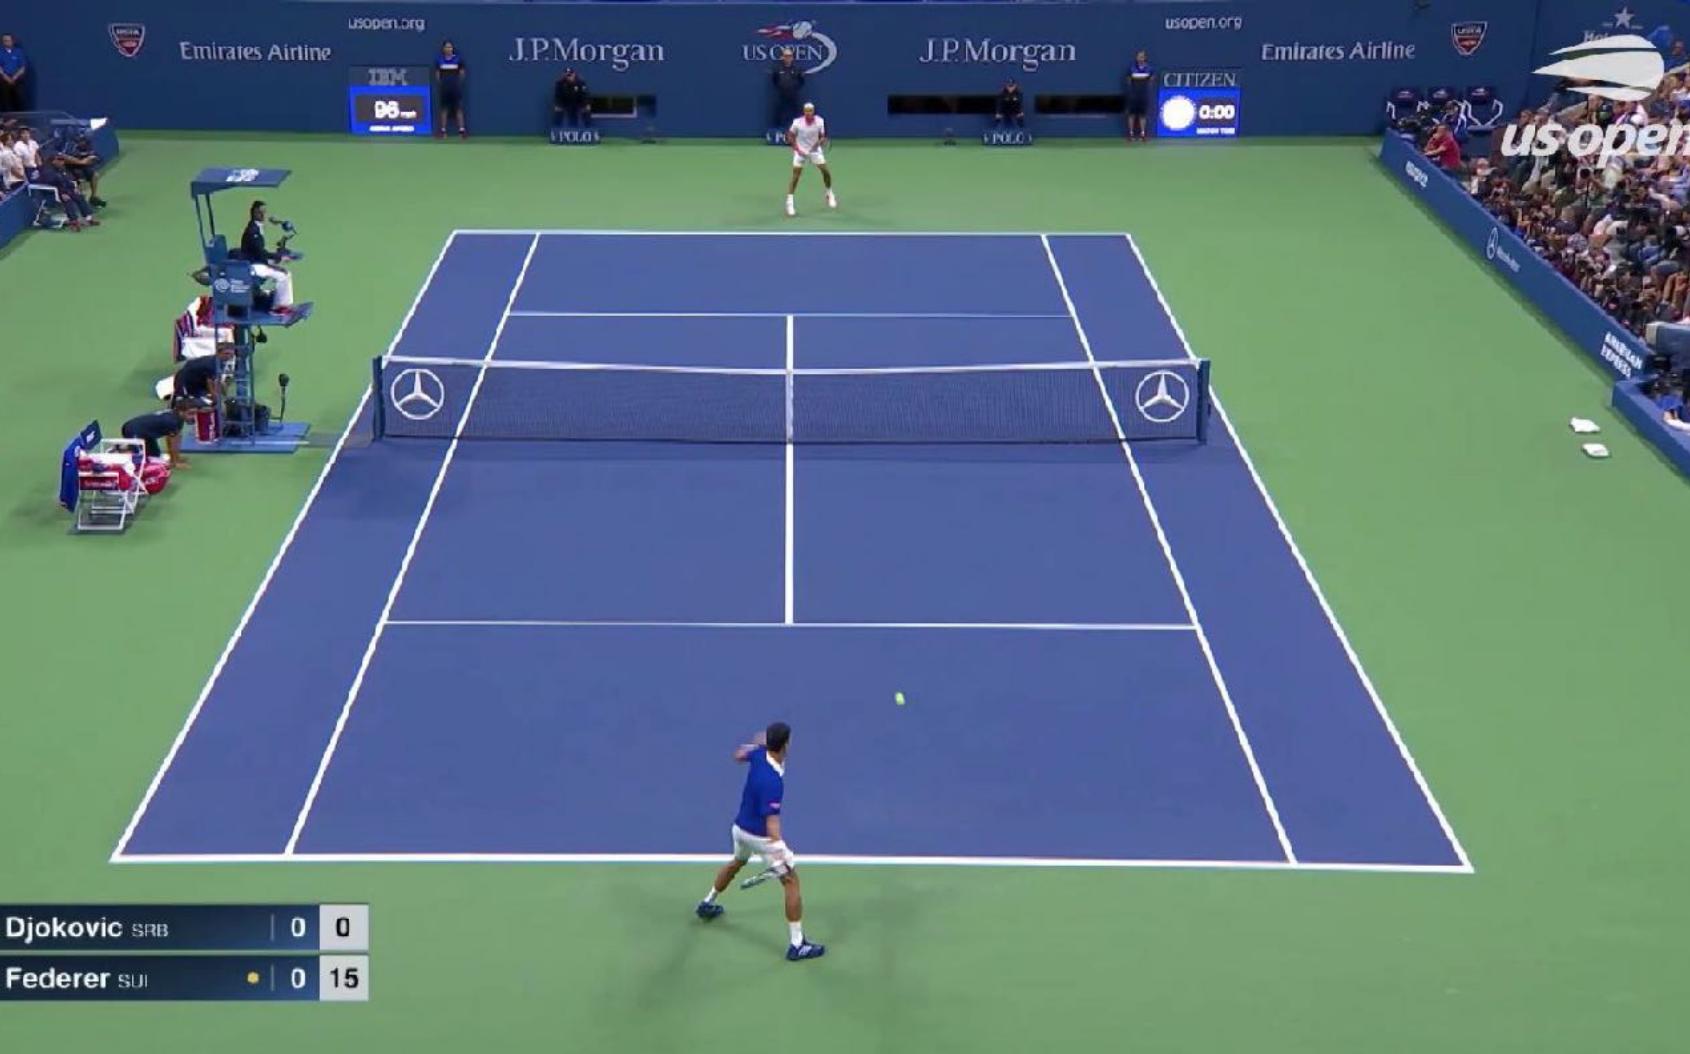
**Compare the description above to the following sequence...**

# Here's an example of that source video

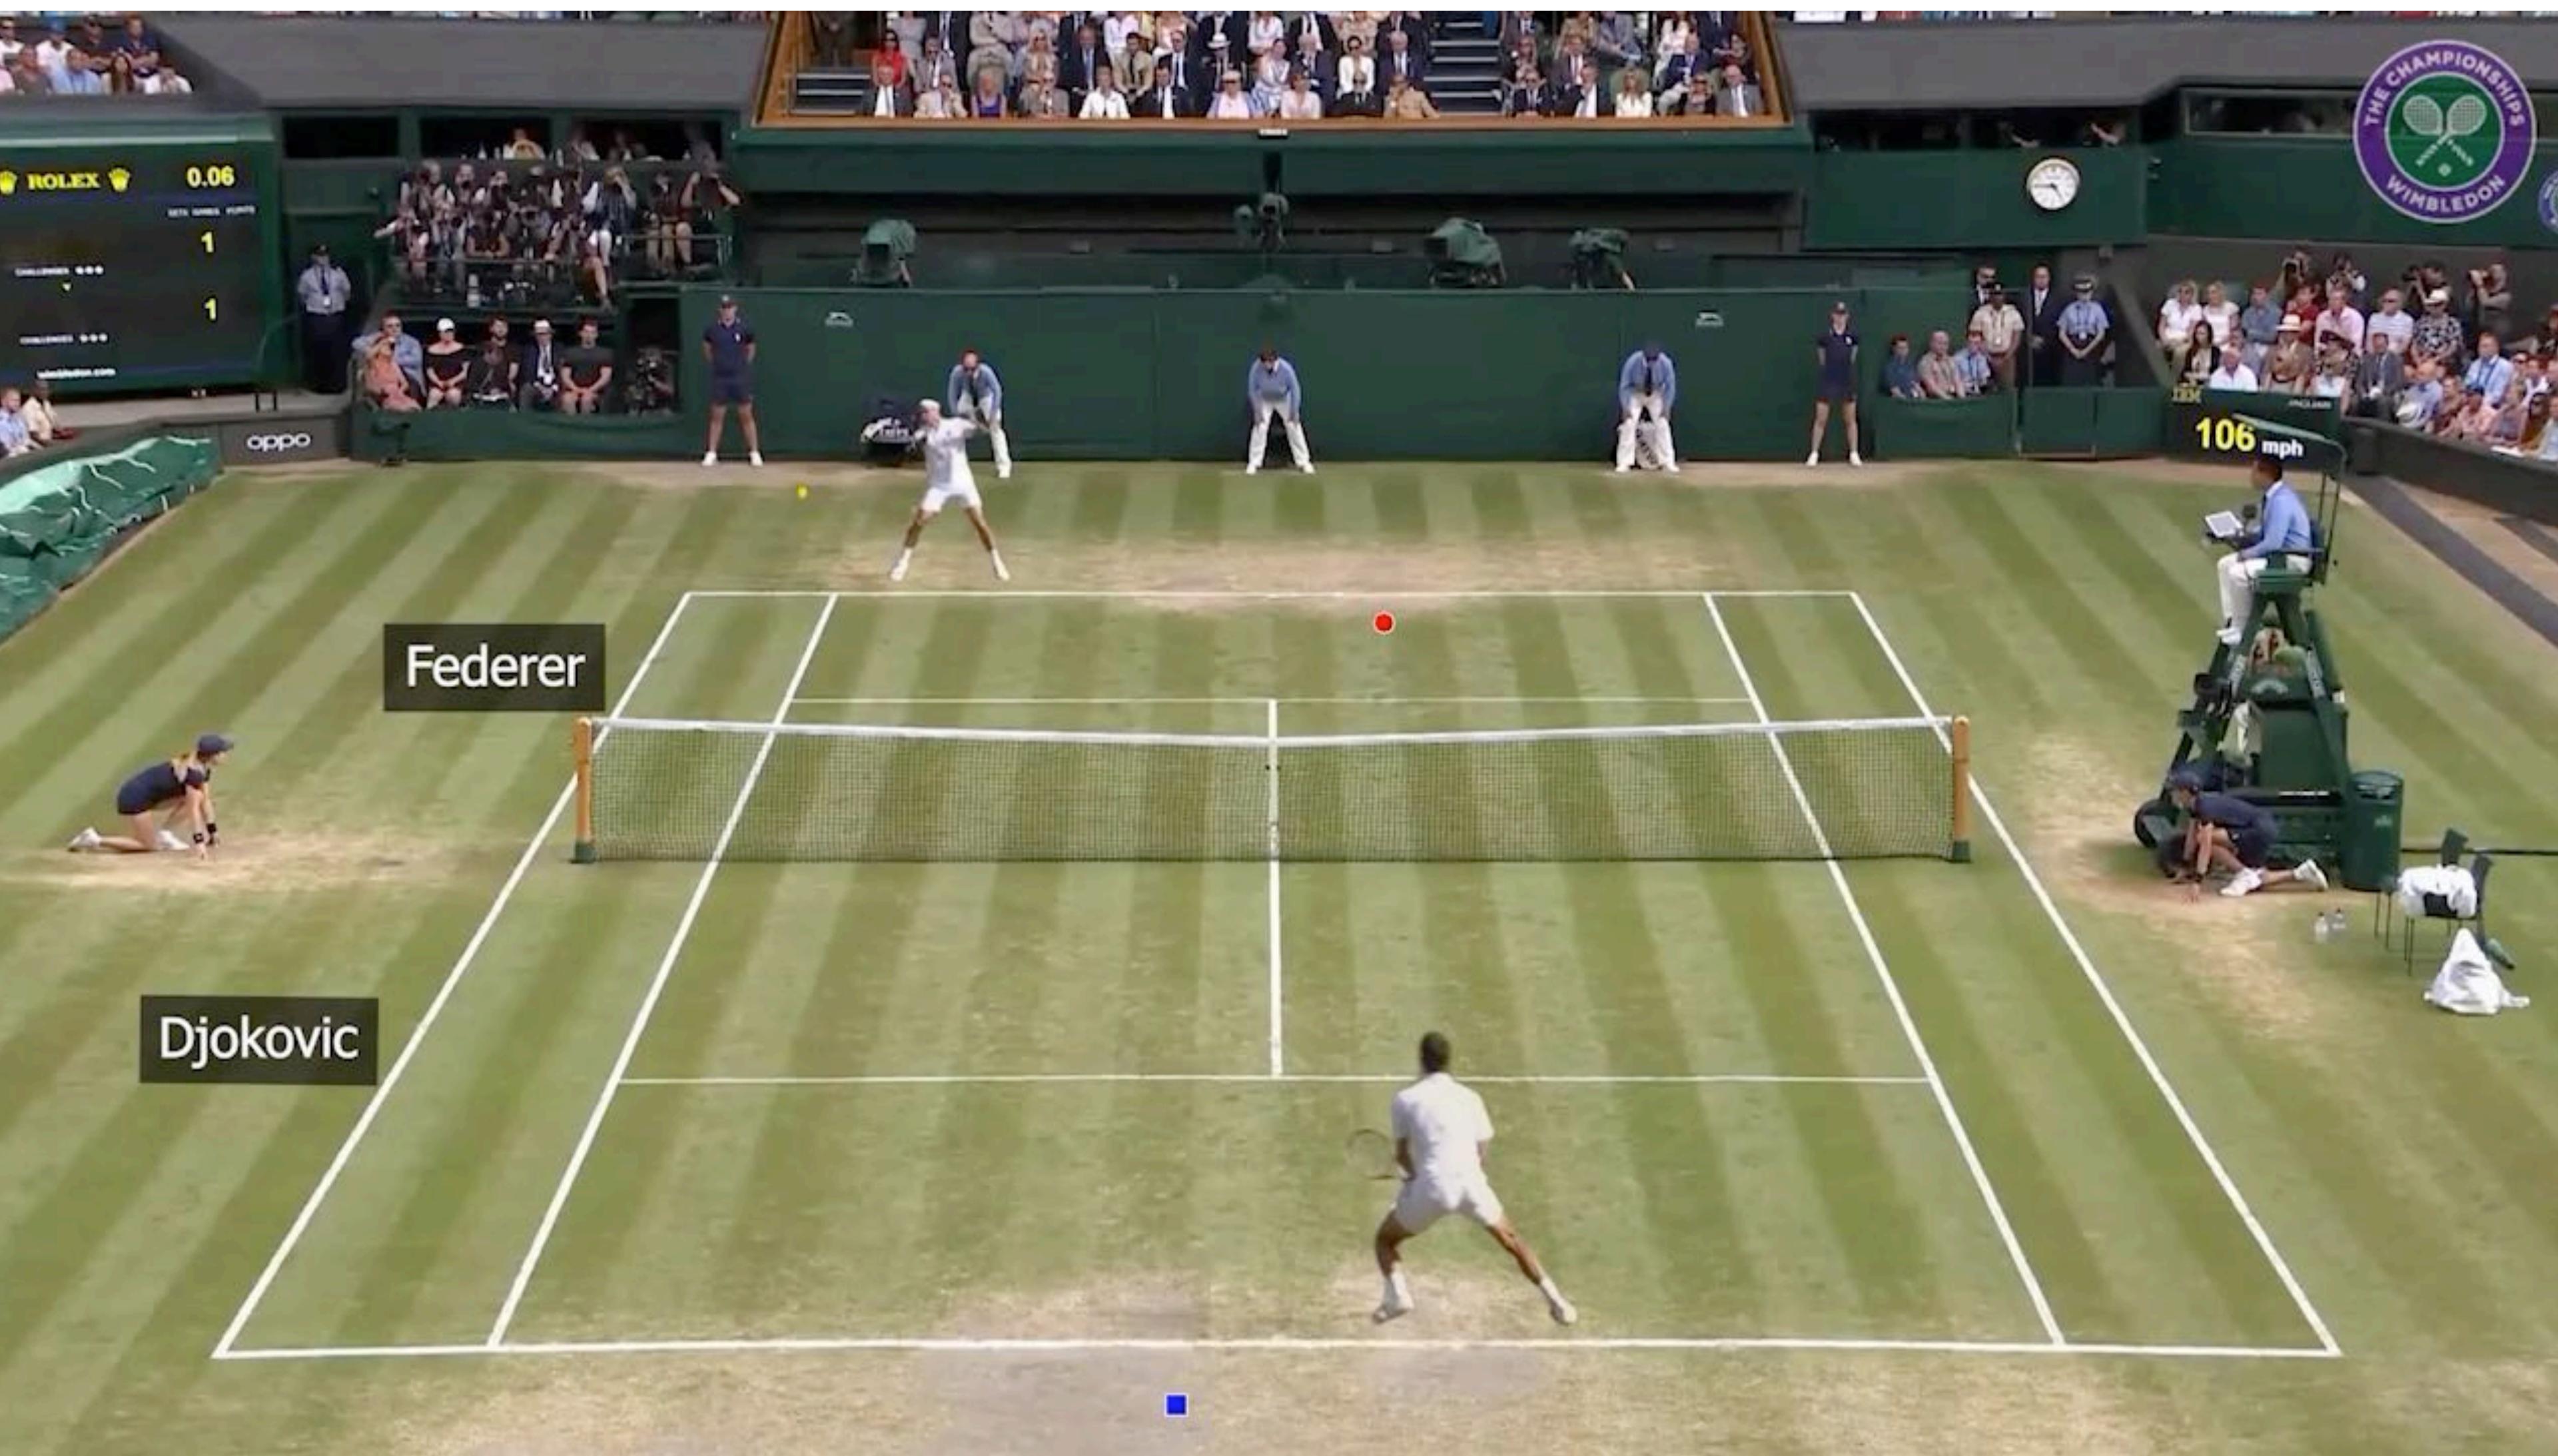


The best way to describe the input data than to just show it!

# And there's a lot of it!



# And here's an example of controllable output



The best way to describe the output we seek is just show the result of the system!

# **Another example:**

- We had a problem in this project: input videos were taken in different lighting conditions, and these lighting differences were the cause of bad results.**
- An anticipated audience question: “what do you mean by lighting differences?”**

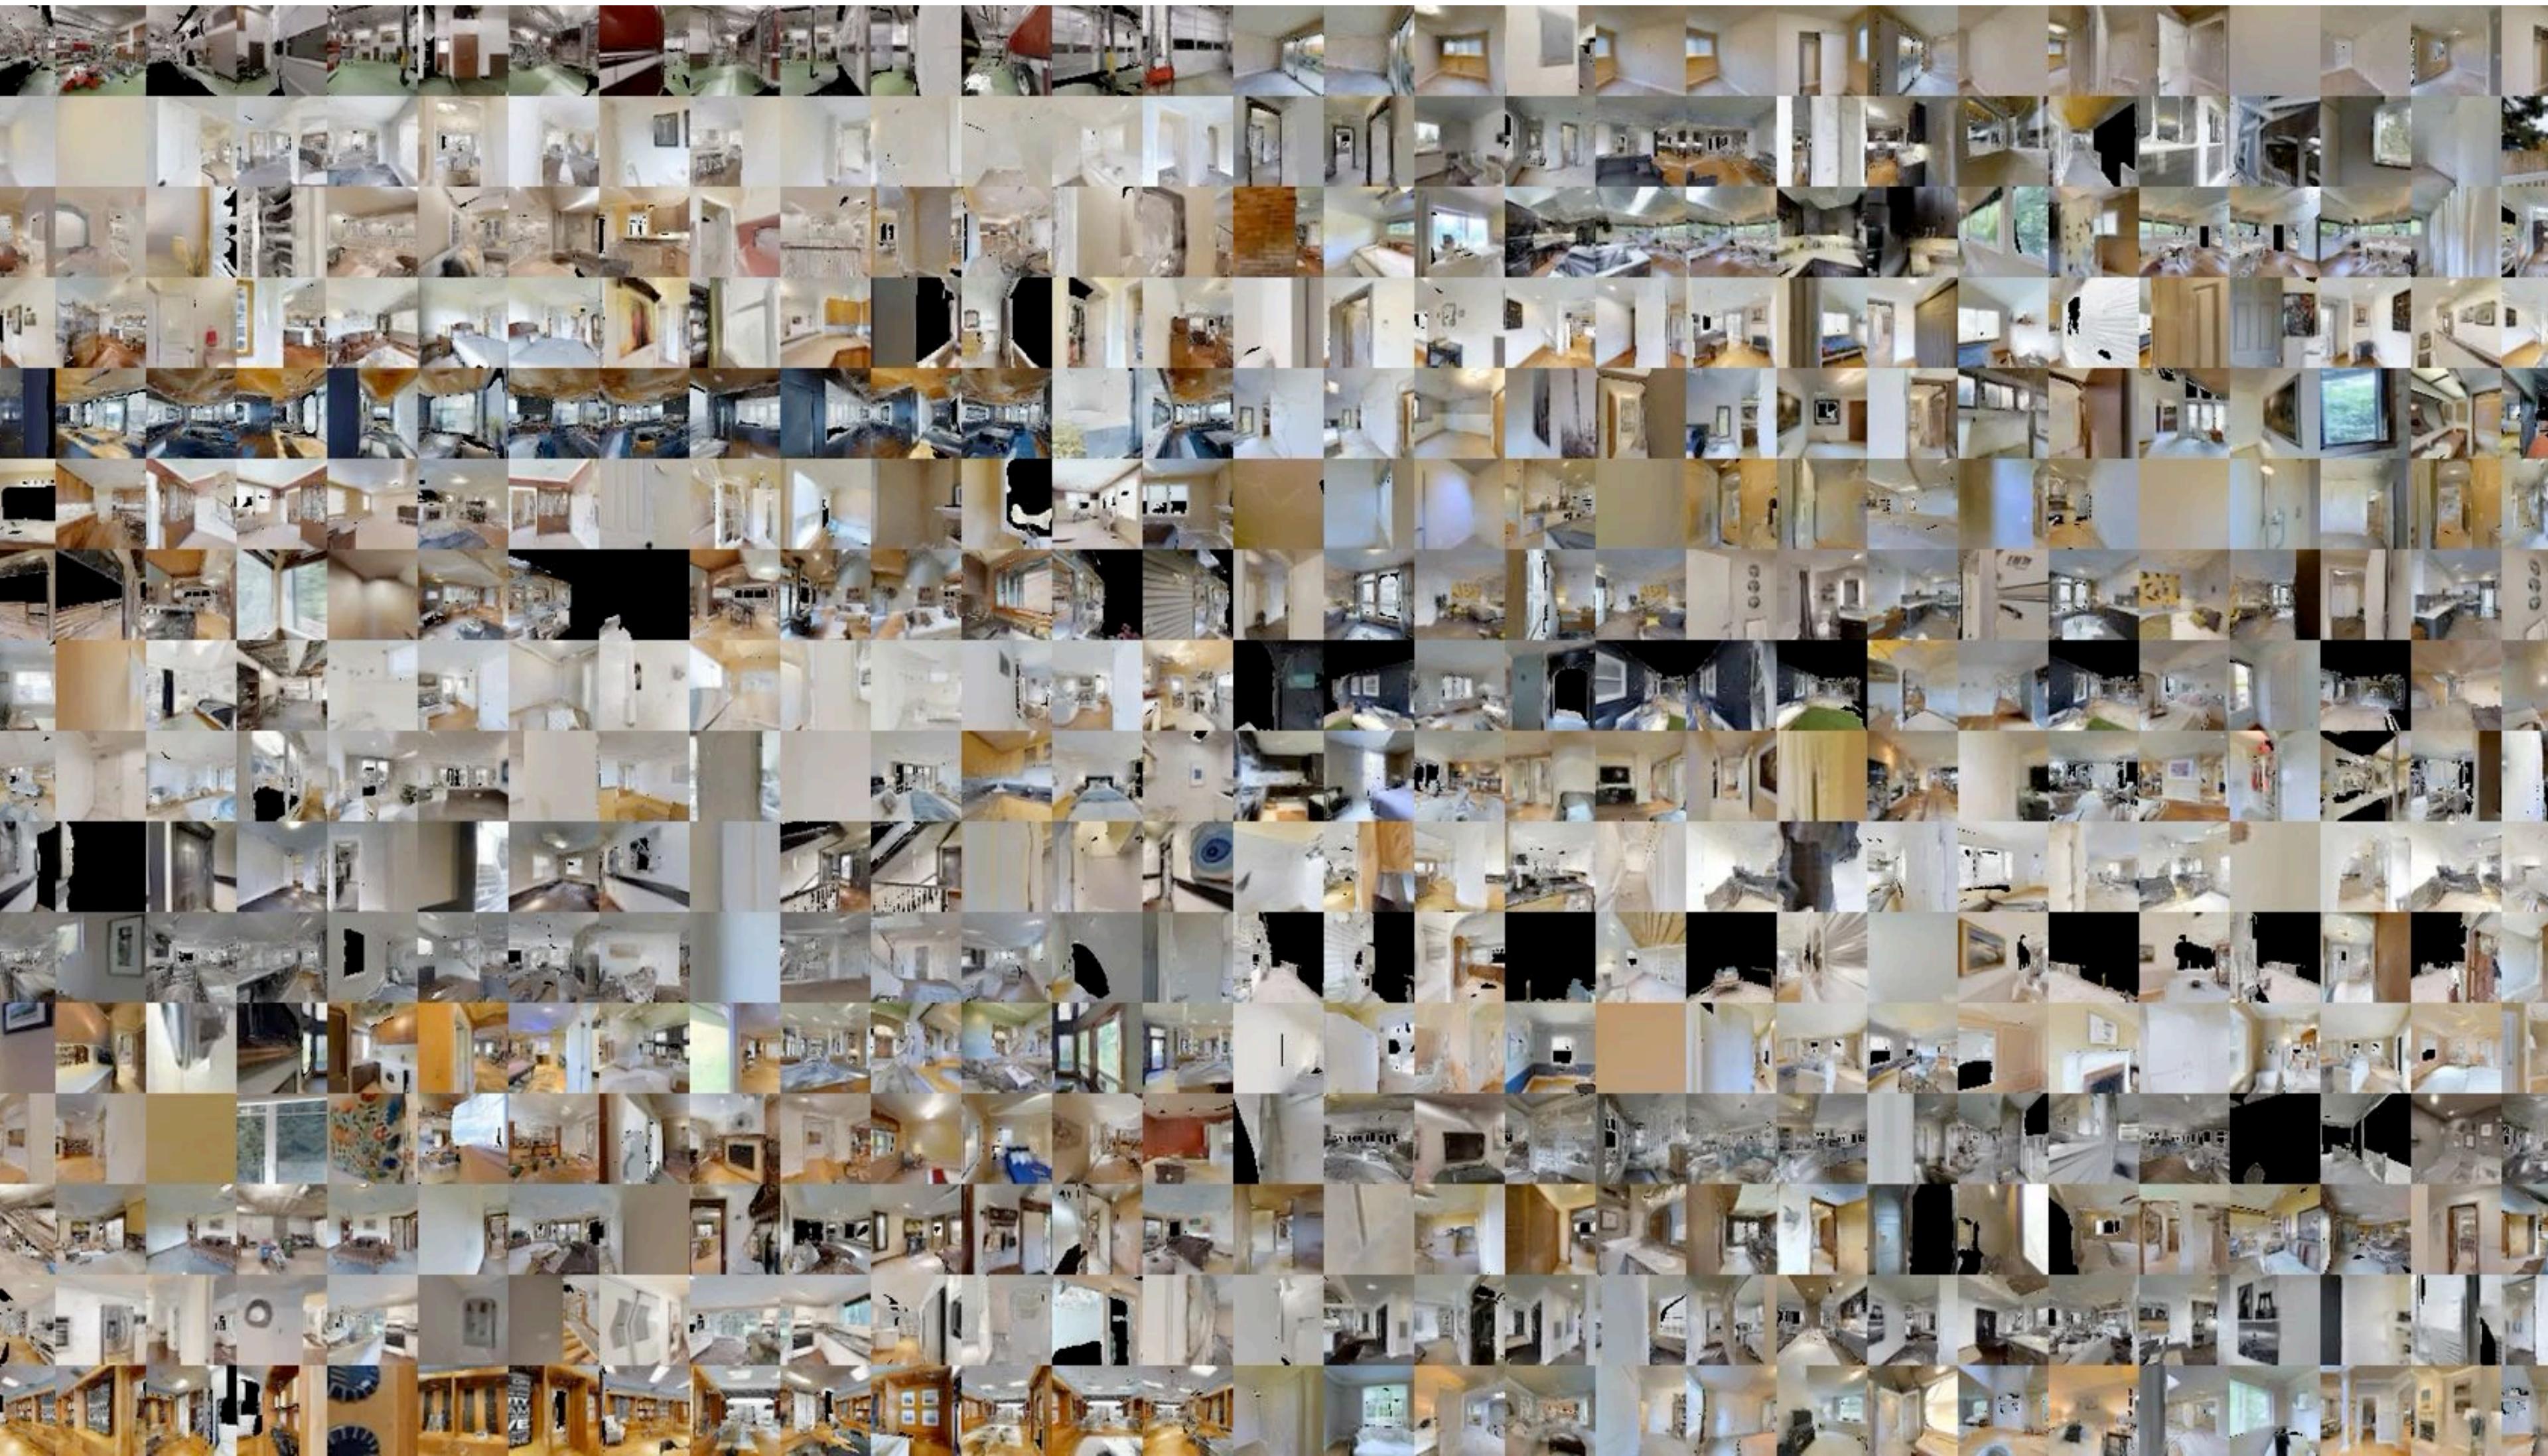
# The problem (lighting differences)



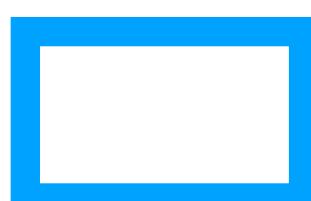
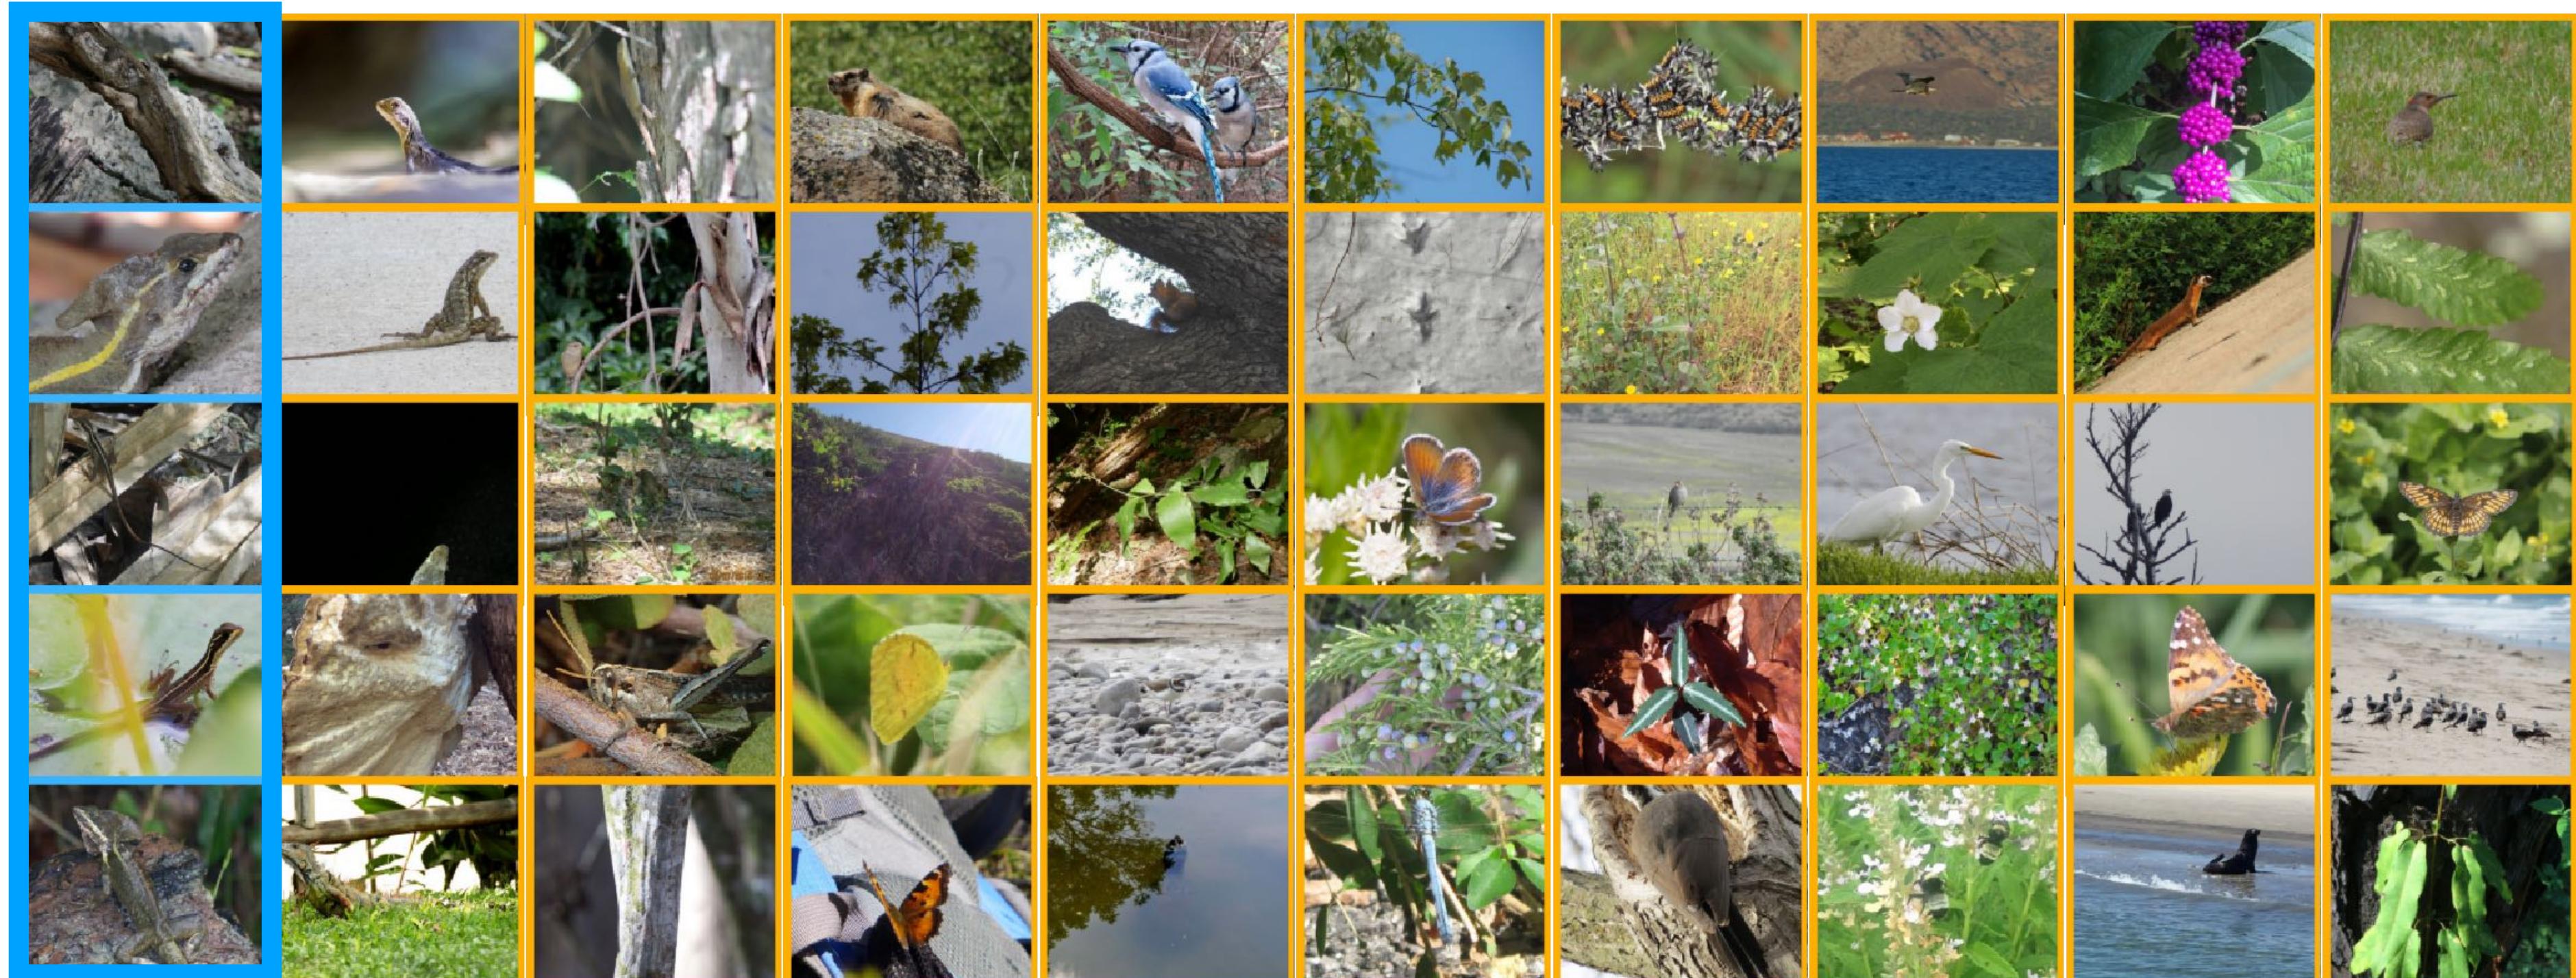
# After the fix



**Another example: we recently created a renderer that achieved high frame-rate rendering by rendering many views of the scene at the same time**



# Another example: it is difficult to train detectors for fine-grained object categories. Consider this...



= positive examples



= negative examples

# **Tip 5**

**The audience prefers not to think (much)**

# An audience has a finite supply of mental effort

- The audience does not want to burn mental effort about things you know and can just tell them.
  - They want to be led by hand through the major steps of your story
  - They do not want to interpret any of your figures or graphs, they want to be directly told how to interpret them (e.g., what to look for in a graph).
  - They want to be told about your key assumptions
- The audience does want to spend their energy thinking about:
  - Potential problems/limitations with what you did (did you consider all edge cases? Is your evaluation sound?)
  - Implications of your approach to their work
  - Connections to their own work

**Which leads me to...**

**The audience does not want to think about  
“why” you are telling them something.**

# Tip 6

**Surprises\* are almost always bad:  
Say where you are going and why you must go  
there before you say what you did.**

\* I am referring to surprises in talk narrative and/or exposition. A surprising result is great.

# Give the why before the what

- Why provides the listener context for...
  - Compartmentalizing: assessing how hard they should pay attention (is this a critical idea, or just an implementation detail?). Especially useful if they are getting lost.
  - Understanding how parts of the talk relate (“Why is the speaker now introducing a new optimization framework?”)
- In the algorithm description:
  - “We need to first establish some terminology”
  - “Even given X, the problem we still haven’t solved is...”
  - “Now that we have defined a cost metric we need a method to minimize it...”
- In the results/evaluation:
  - Speaker: “Key questions to ask about our approach are...”
  - Audience: “Thanks! I agree, those are good questions. Let’s see what the results say!”

Two key questions:

- How much does SRDH improve traversal cost when perfect information about shadow rays is present?
- How does the benefit of the SRDH decrease as less shadow ray information is known a priori? (Is a practical implementation possible?)

# Big surprises in a narrative are a bad sign

- Ideally, you want the audience to always be able to anticipate\* what you are about to say
  - This means: your story is so clear it's obvious!
  - It also means the talk is really easy to present without notes or text on slides (it just flows)
- If you are practicing your talk, and you keep forgetting what's coming on the next slide (that is, you can't anticipate it)...
  - This means: you probably need to restructure your talk because a clear narrative is not there.
  - It's not even obvious to you! Ouch!

\* Credit to Abhinav Gupta for suggesting the term anticipation, and for the example on this slide

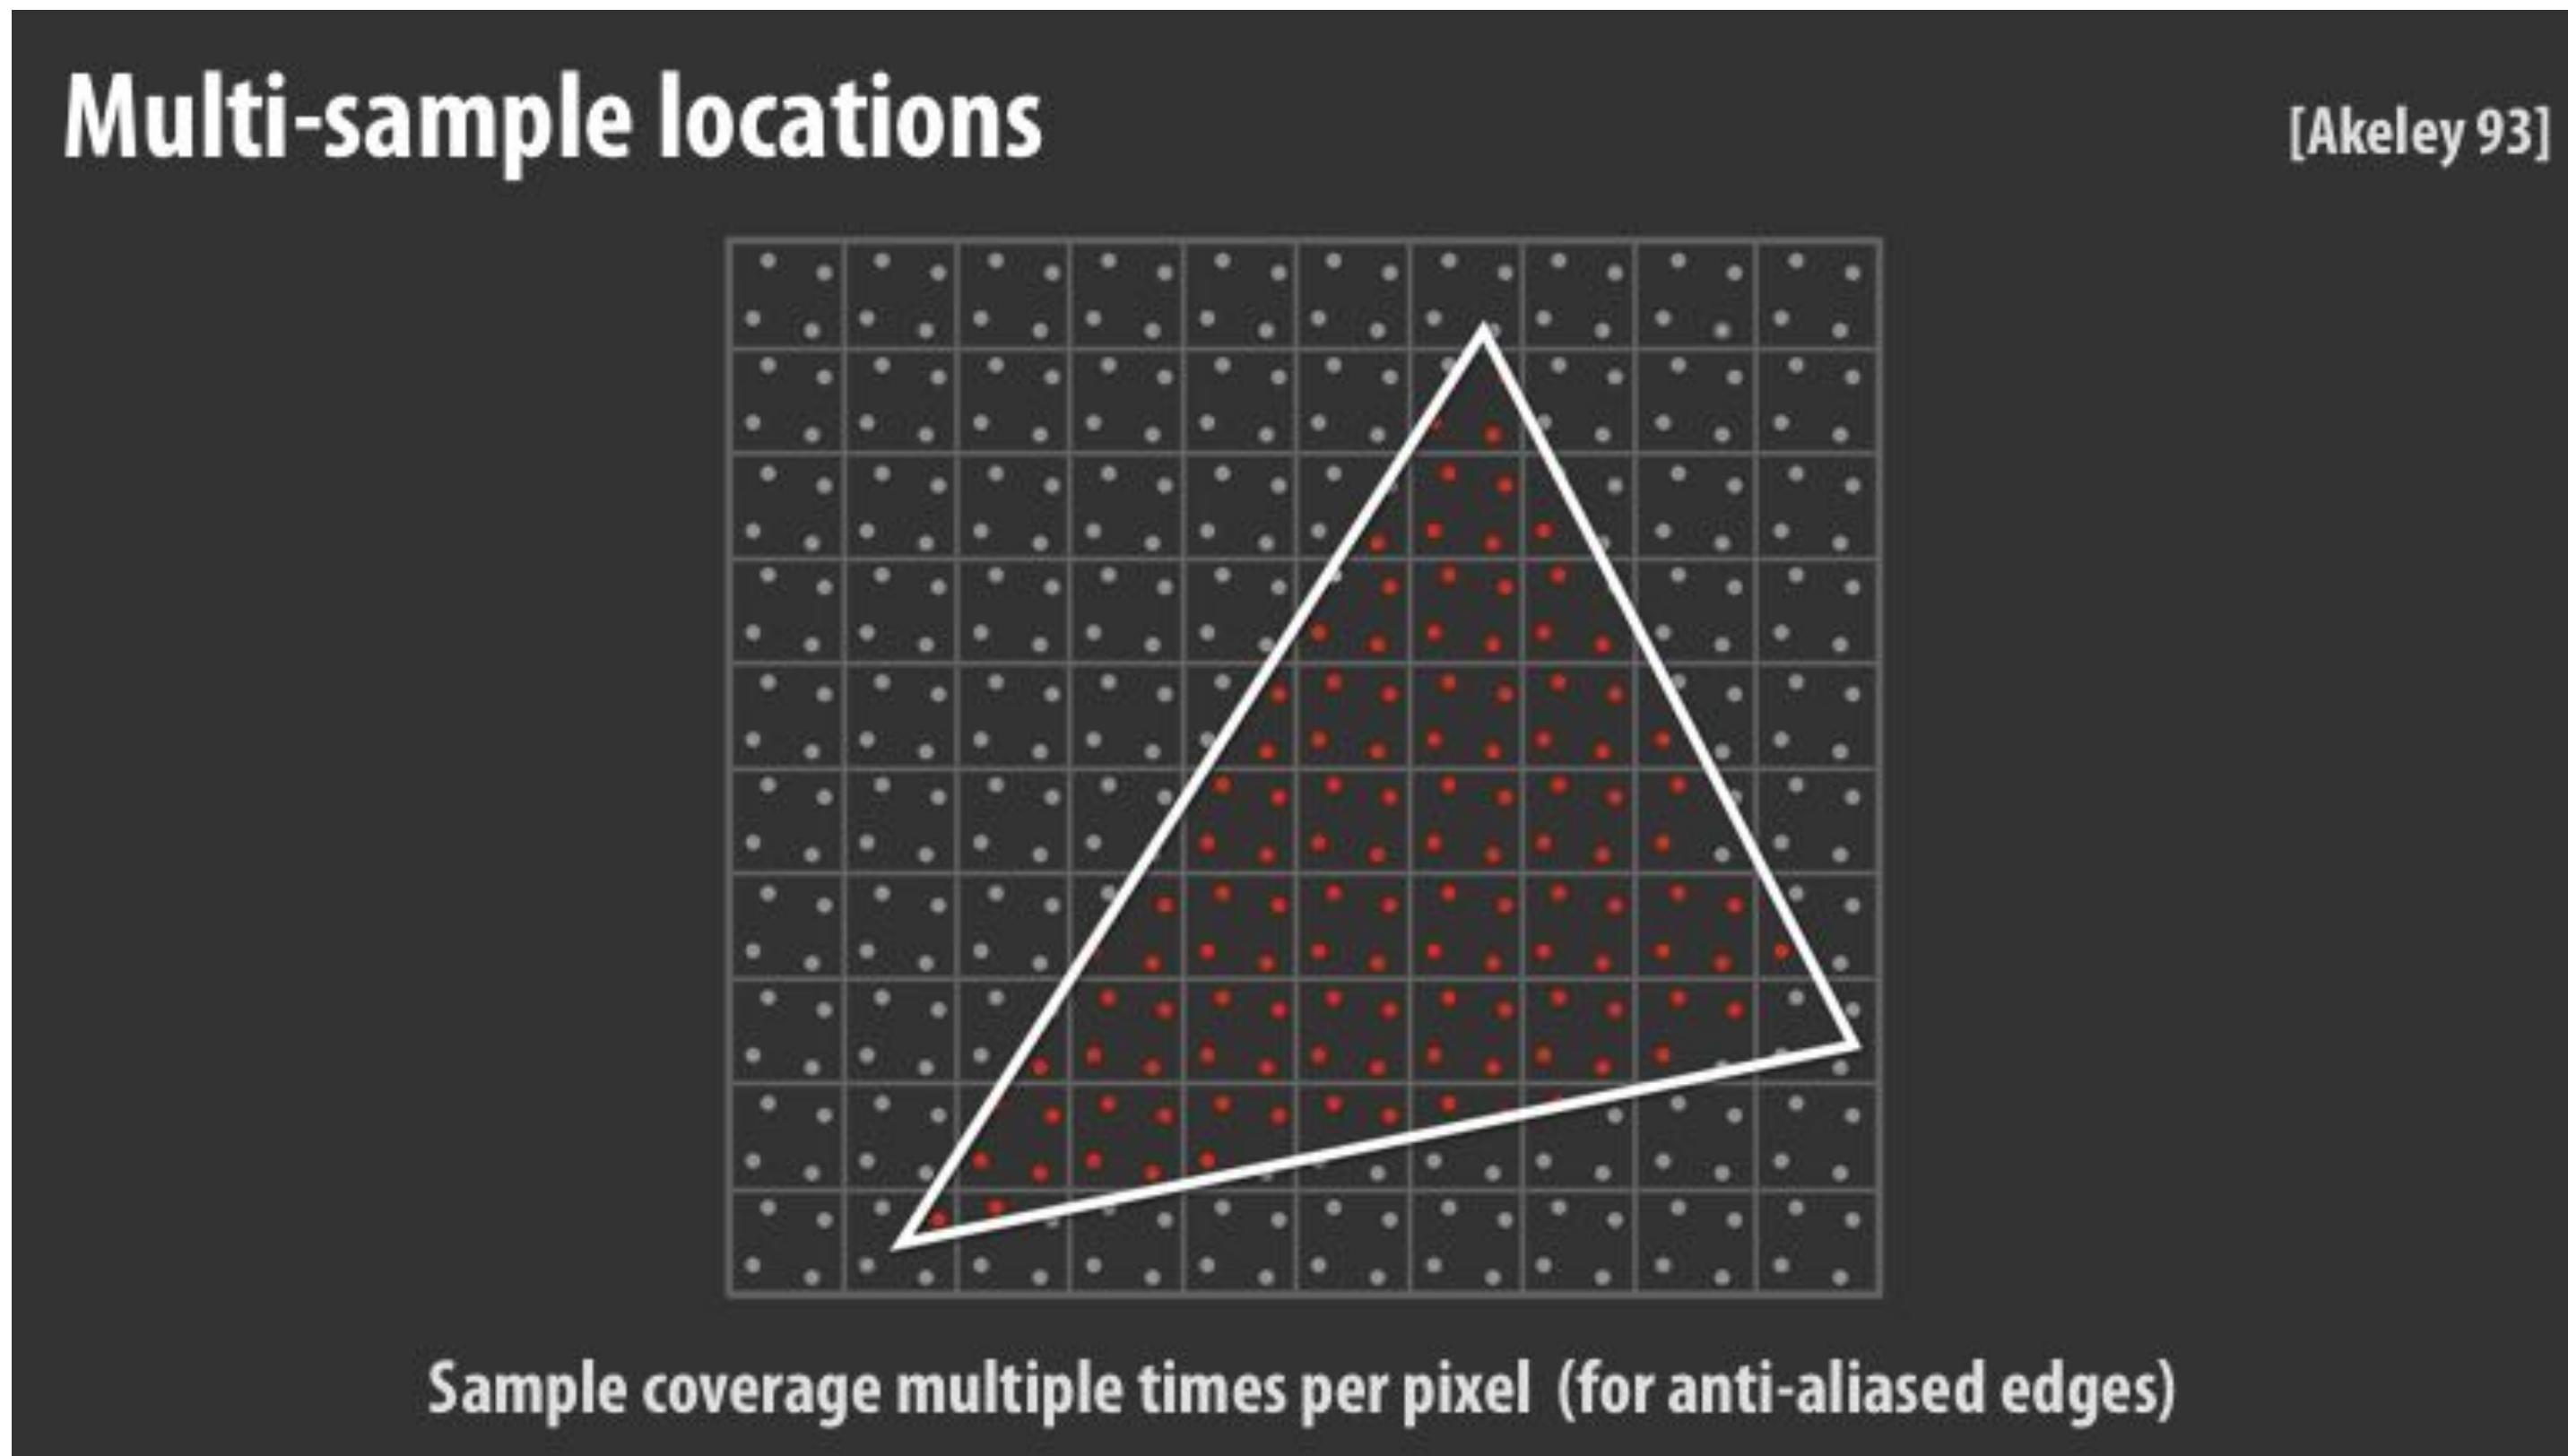
# Tip 7

**Always, always, always  
explain any figure or graph**

**(remember, the audience does not want to think about things you can tell them)**

# Explain every figure

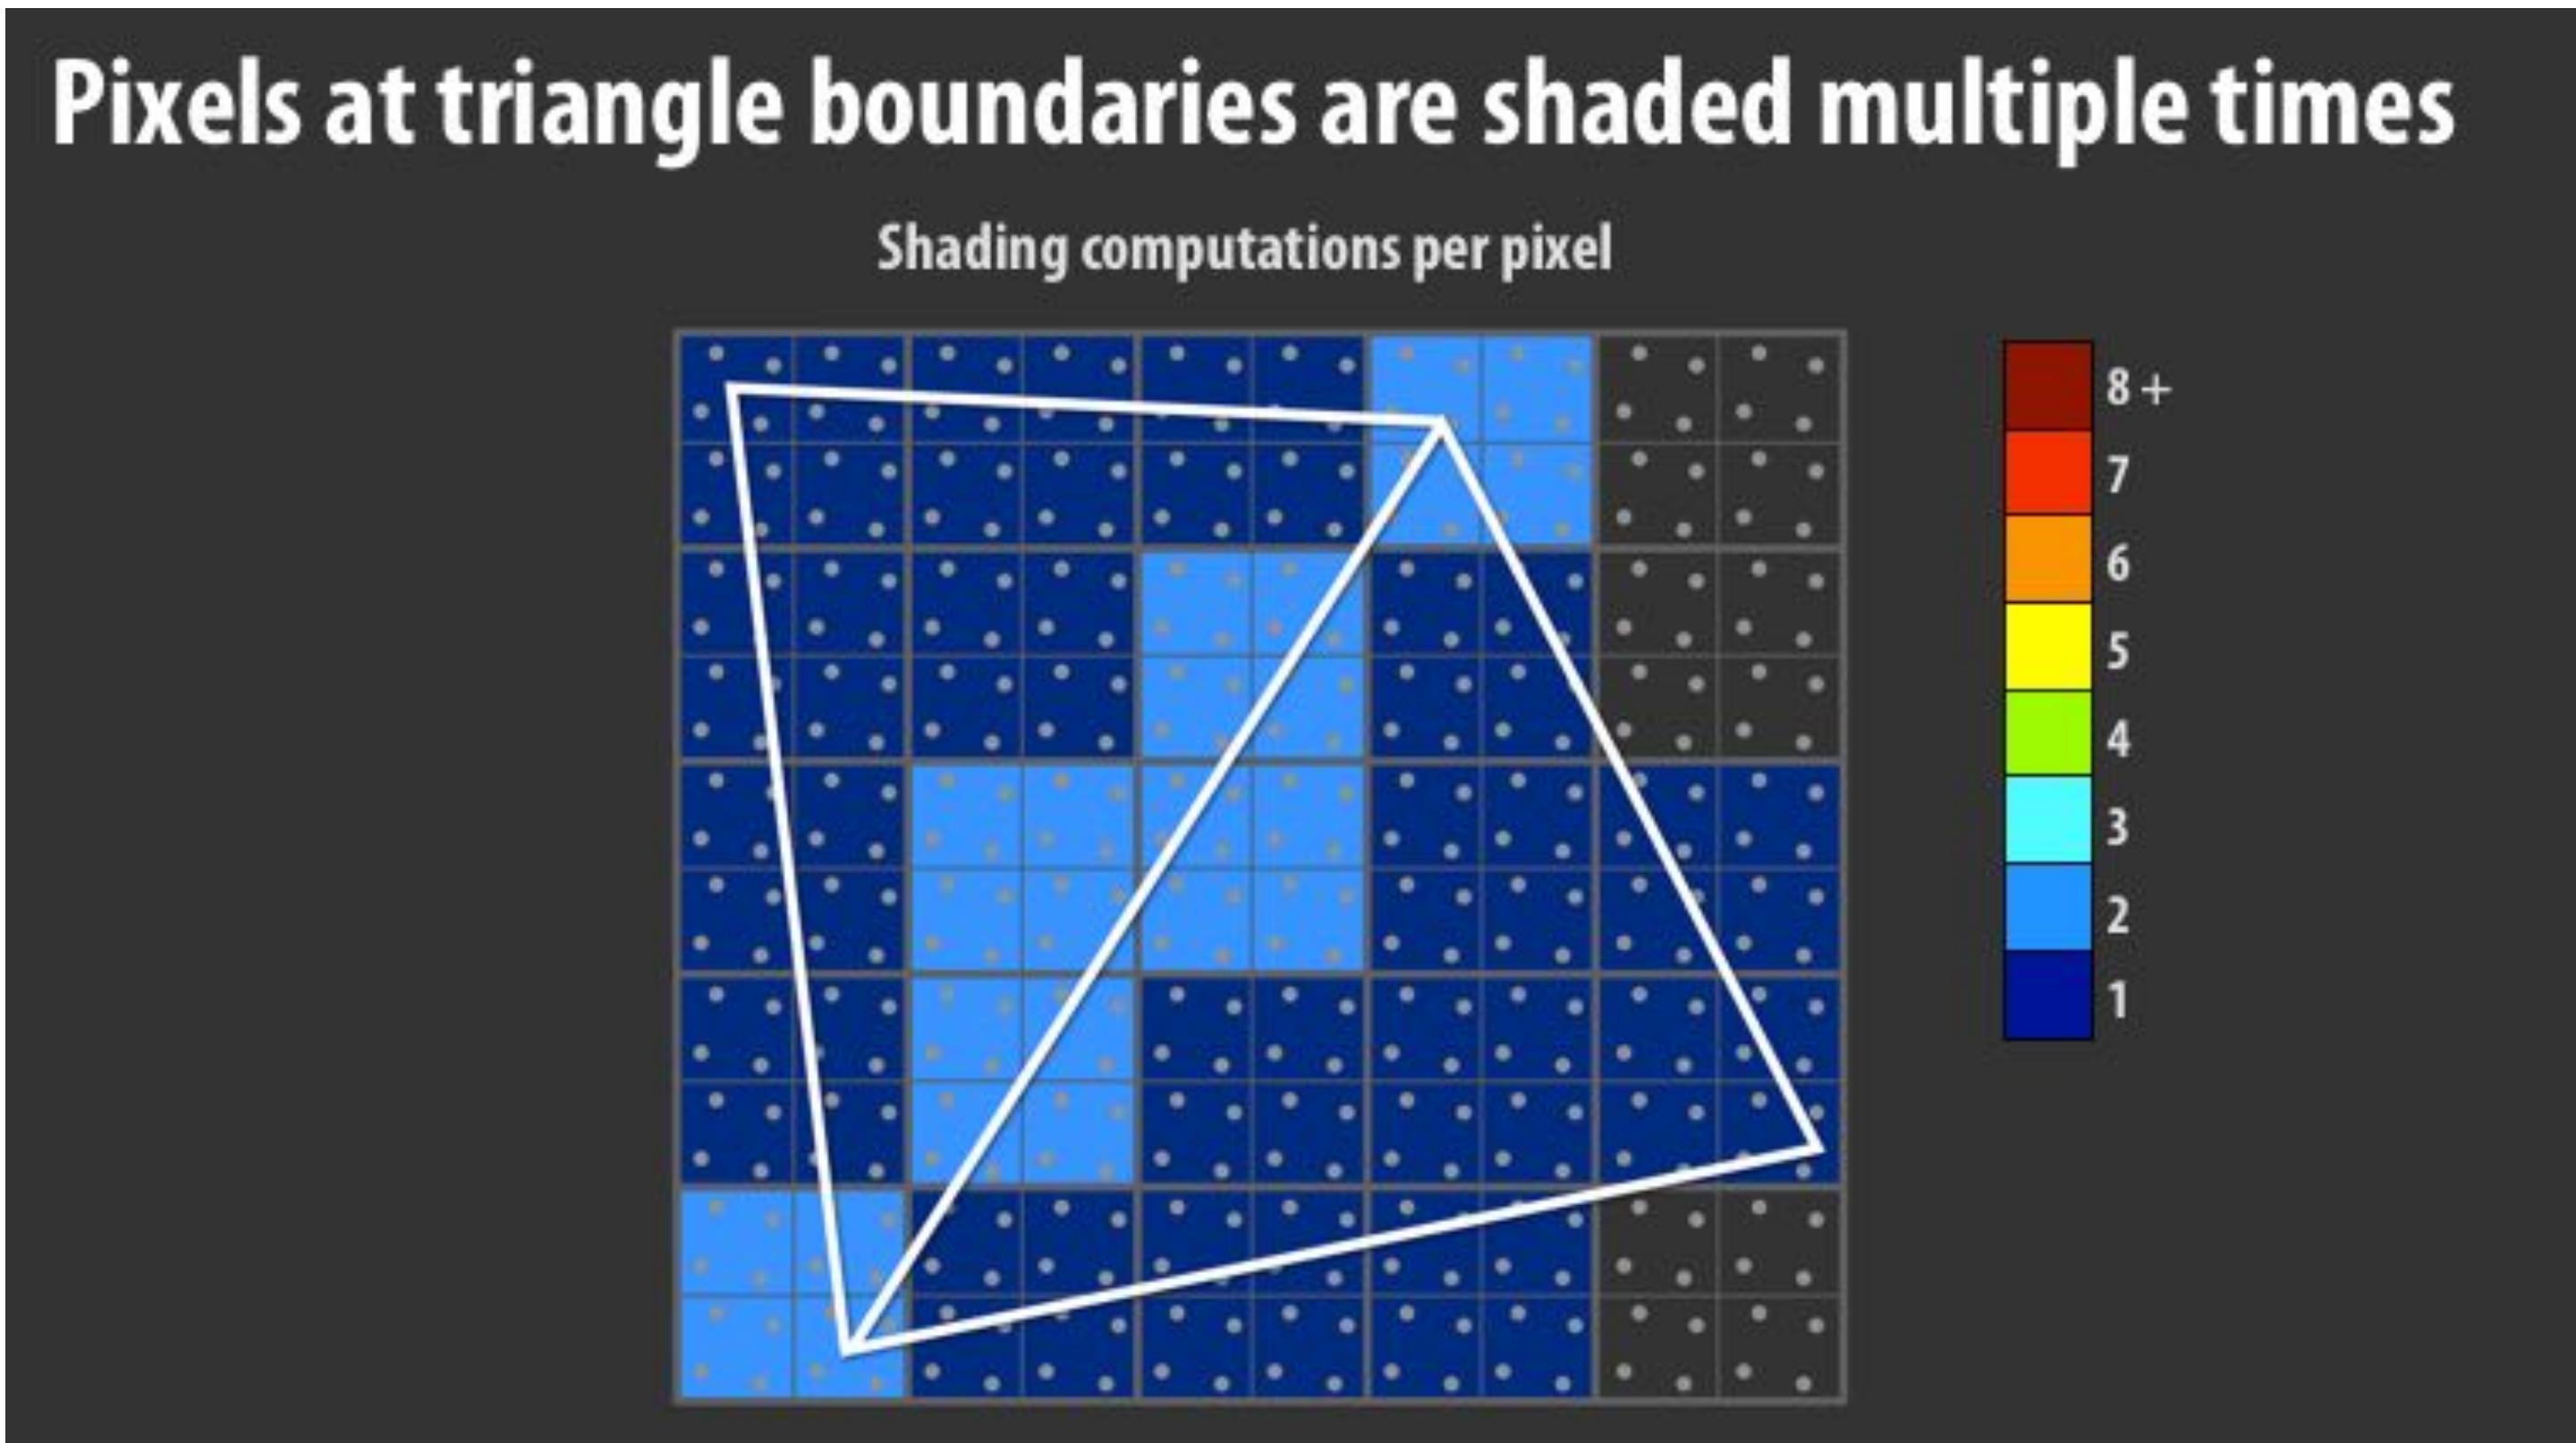
- Explain every visual element in the figure (never make the audience decode a figure)
- Refer to highlight colors explicitly (explain why the visual element is highlighted)



Example voice over: "Here I'm showing you a pixel grid, a projected triangle, and the location of four sample points at each pixel. Sample points falling within the triangle are colored red."

# Explain every figure

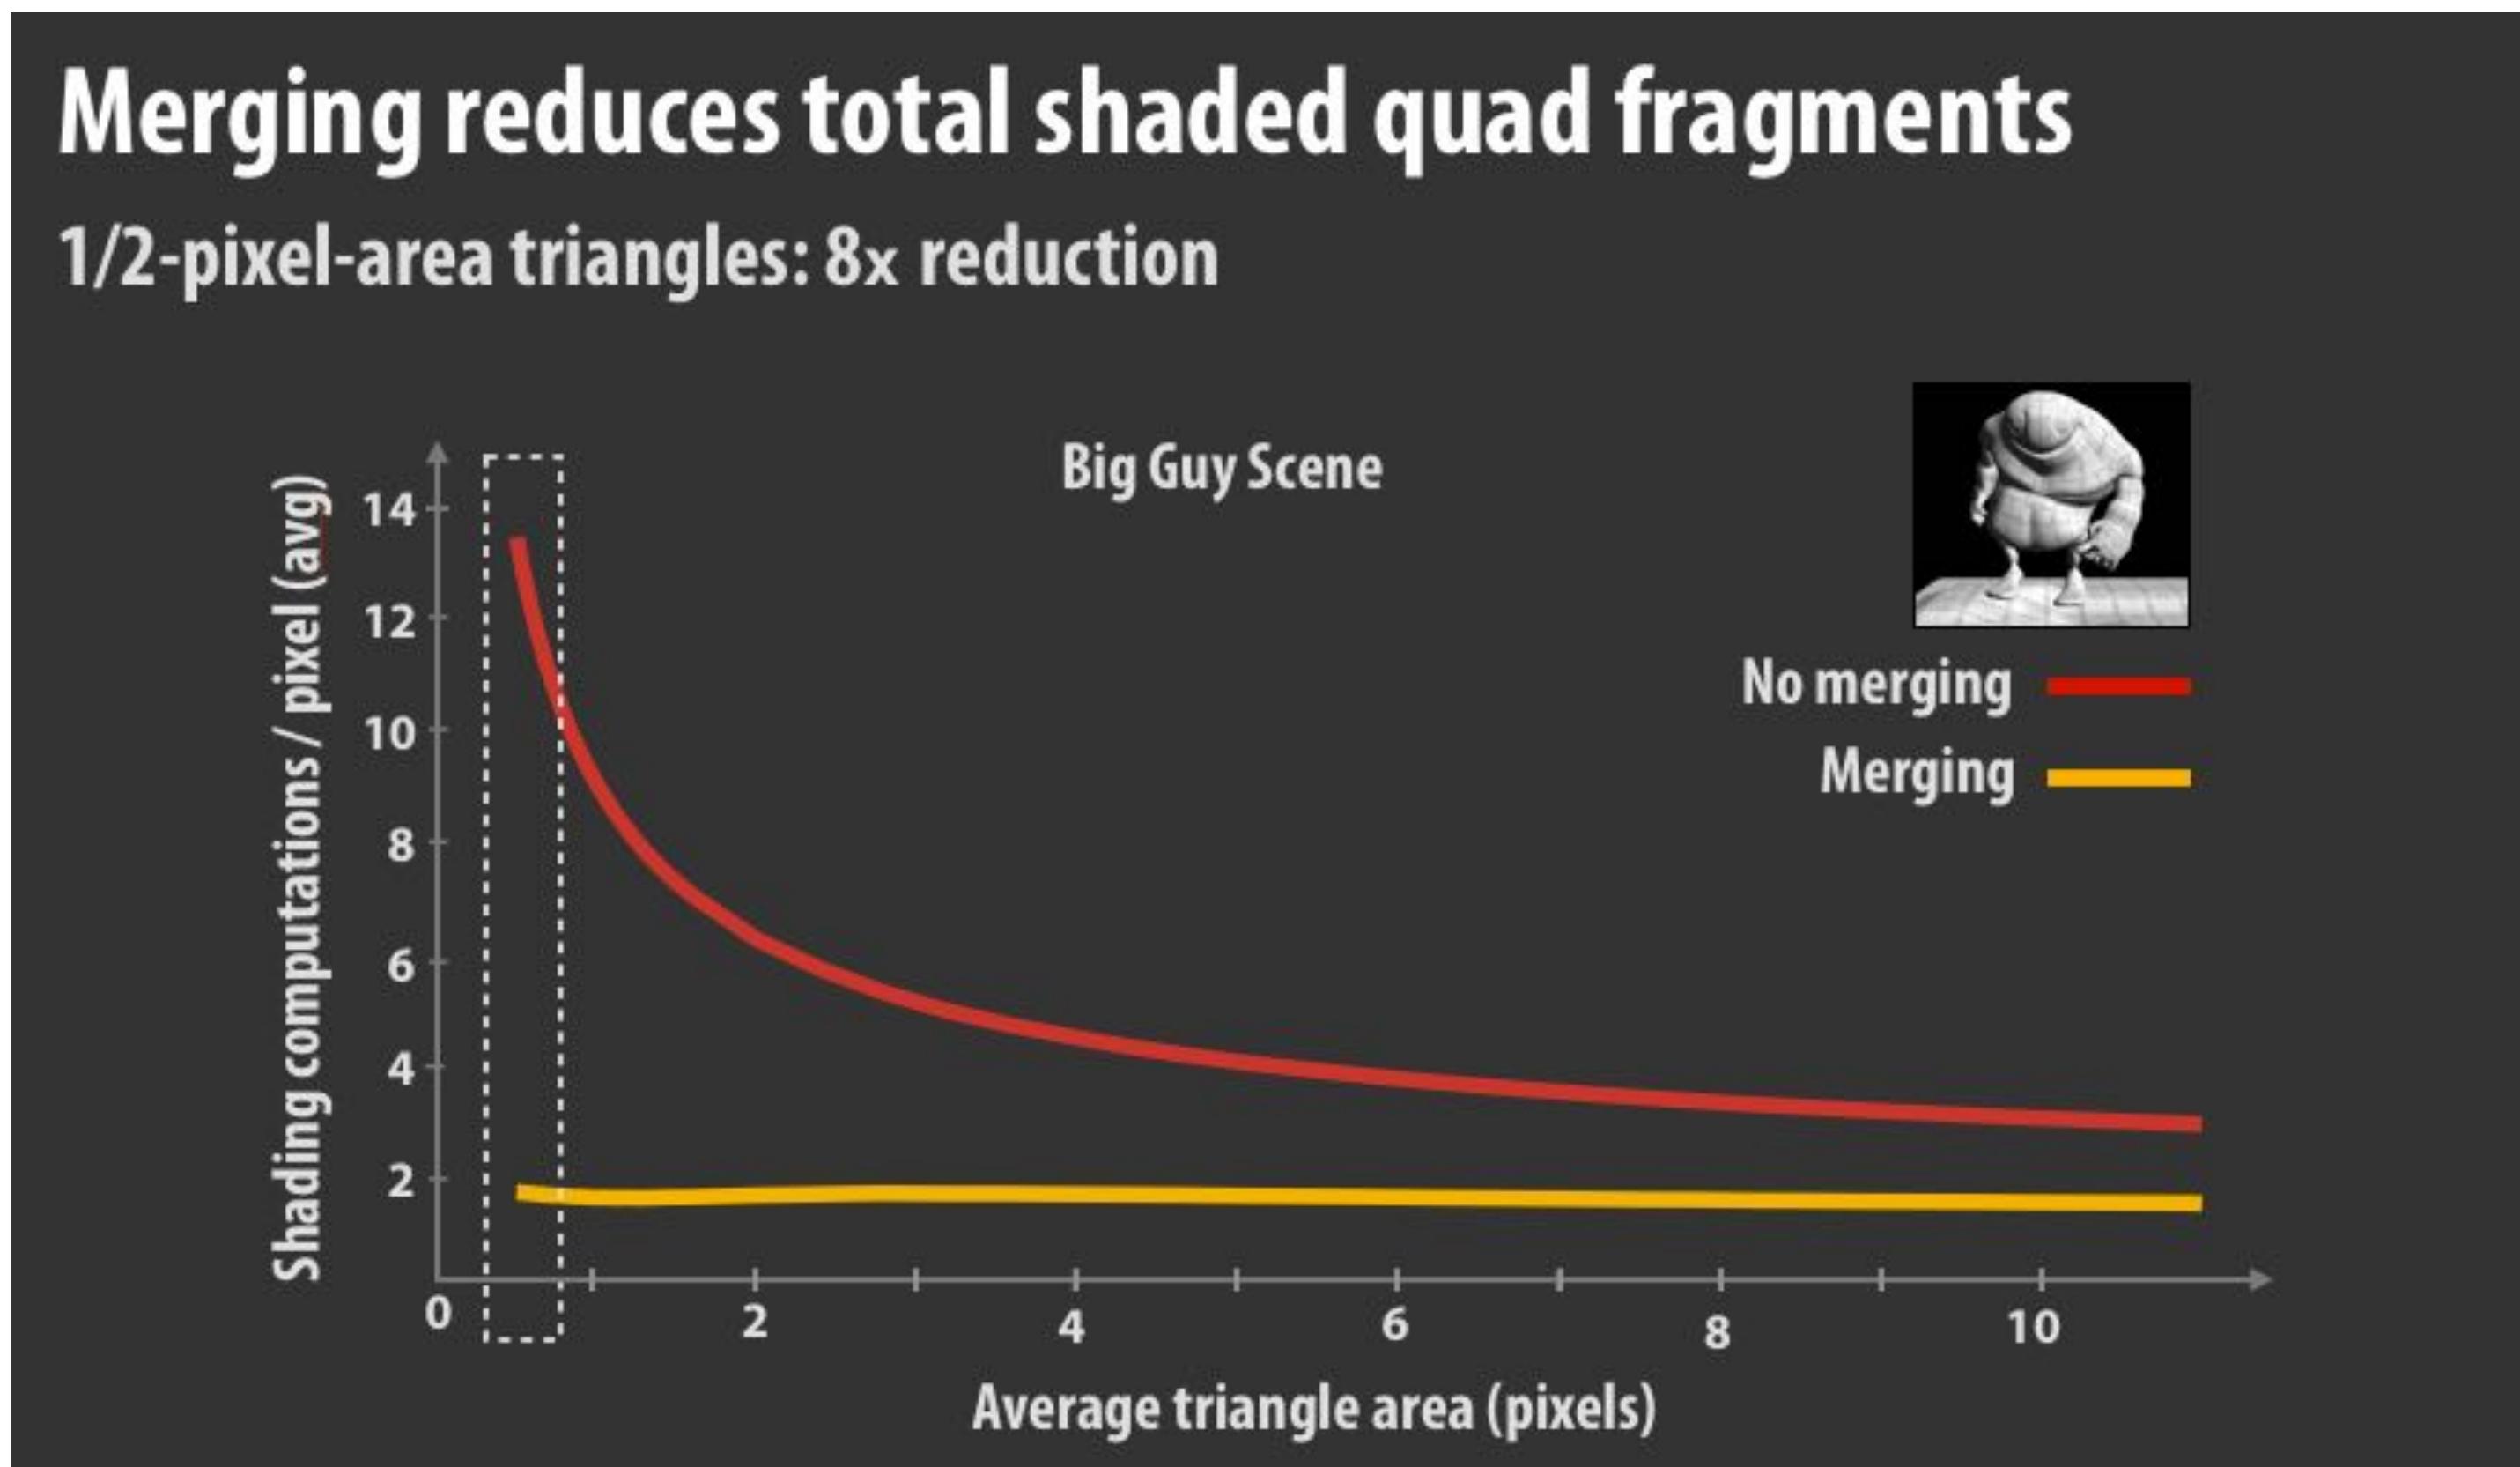
- Lead the listener through the key points of the figure
- Useful phrase: “As you can see...”
  - It’s like verbal eye contact. It keeps the listener engaged and makes the listener happy... “Oh yeah, I can see that! I am following this talk!”



Example voice over: “Now I’m showing you two adjacent triangles, and I’m coloring pixels according to the number of shading computations that occur at each pixel as a result of rendering these two triangles. As you can see from the light blue region, pixels near the boundary of the two triangles get shaded twice.

# Explain every results graph

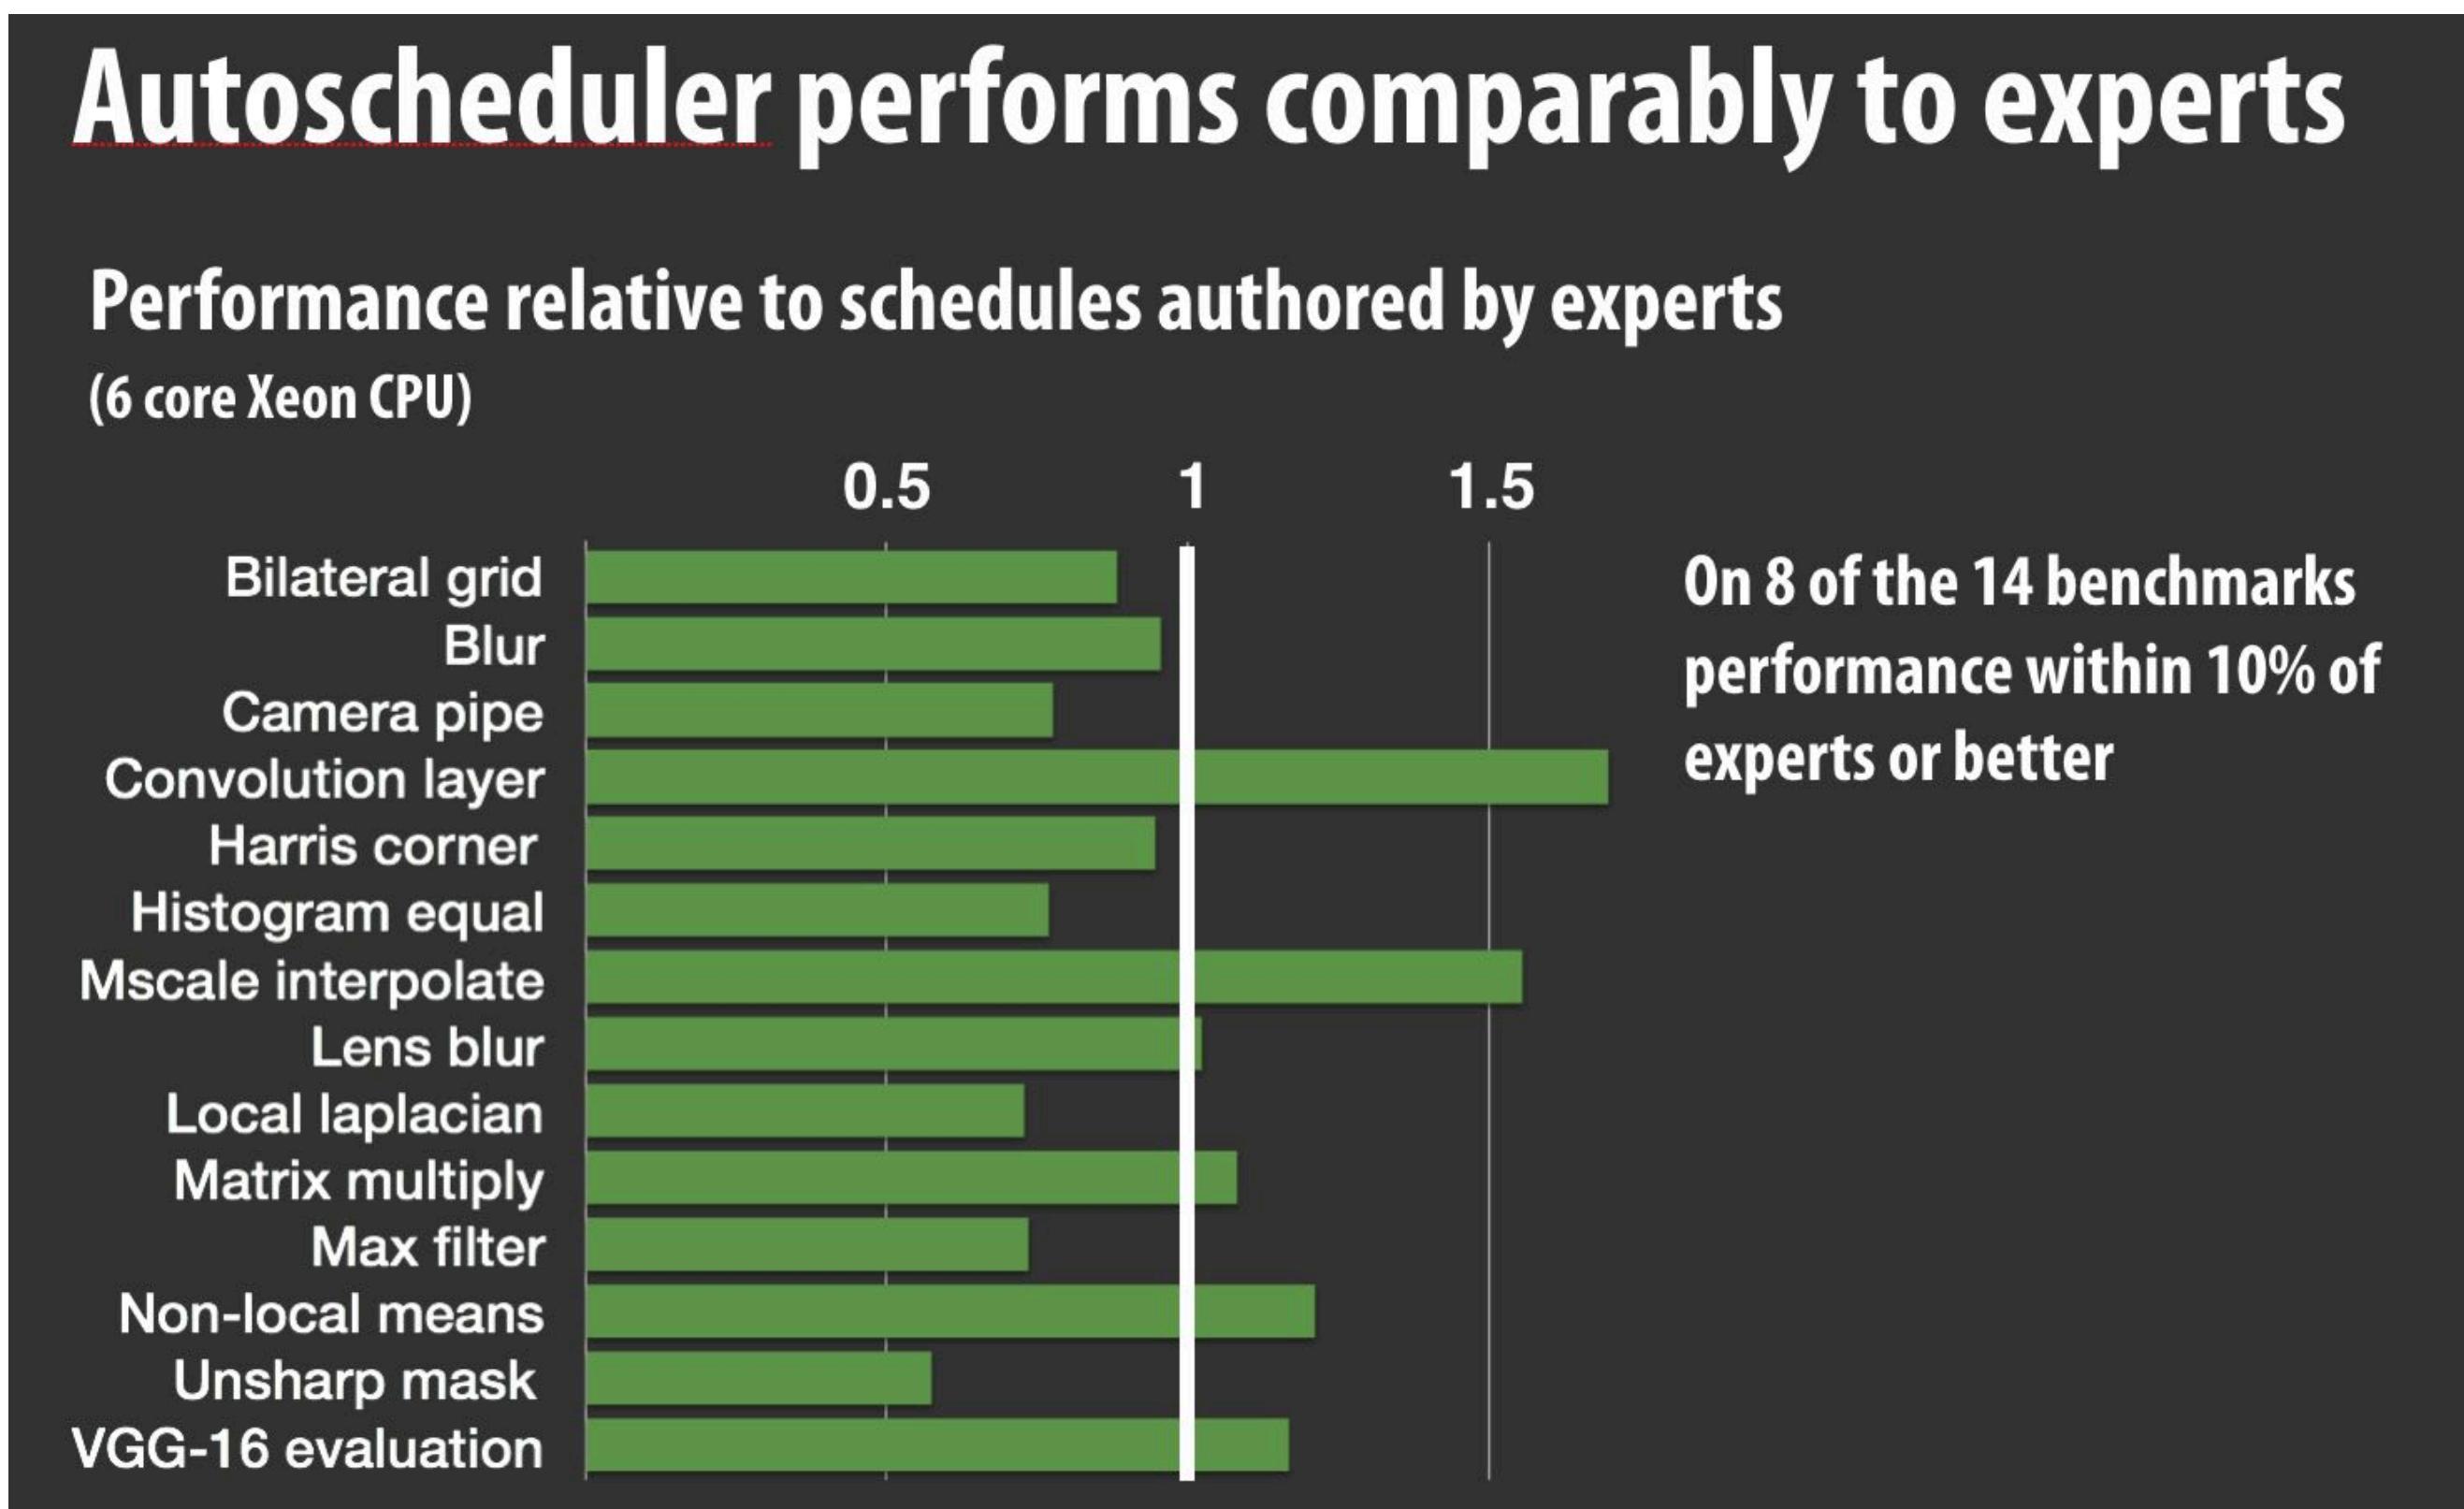
- May start with a general intro of what the graph will address (“anticipate” the result)
- Then describe the axes (and your axes better have labels!)
- Then describe the one point that you wish to make with this results slide (more on this later!)



Example voice over: “Our first questions were about performance: how much did the algorithm reduce the number of the shading computations? And we found out that the answer is a lot. This figure plots the number of shading computations per pixel when rendering different tessellations of the big guy scene. X-axis gives triangle size. If you look at the left side of the graph, which corresponds to a high-resolution micropolygon mesh, you can see that merging, shown by yellow line, shades over eight times less than the convention pipeline.”

# Explain every results graph

- May start with a general intro of what the graph will address.
- Then describe the axes (your axes better have labels!)
- Then describe the one point that you wish to make with this results slide (more on this later!)



Example voice over: "Our first question was about performance: how fast is the auto scheduler compared to experts? And we found out that it's quite good. This figure plots the performance of the autoscheduler compared to that of expert code. So expert code is 1. Faster code is to the right. As you can see, the auto scheduler is within 10% of the performance of the experts in many cases, and always within a factor of 2."

# **Tip 8**

**In the results section:**

**One point per slide!**

**One point per slide!**

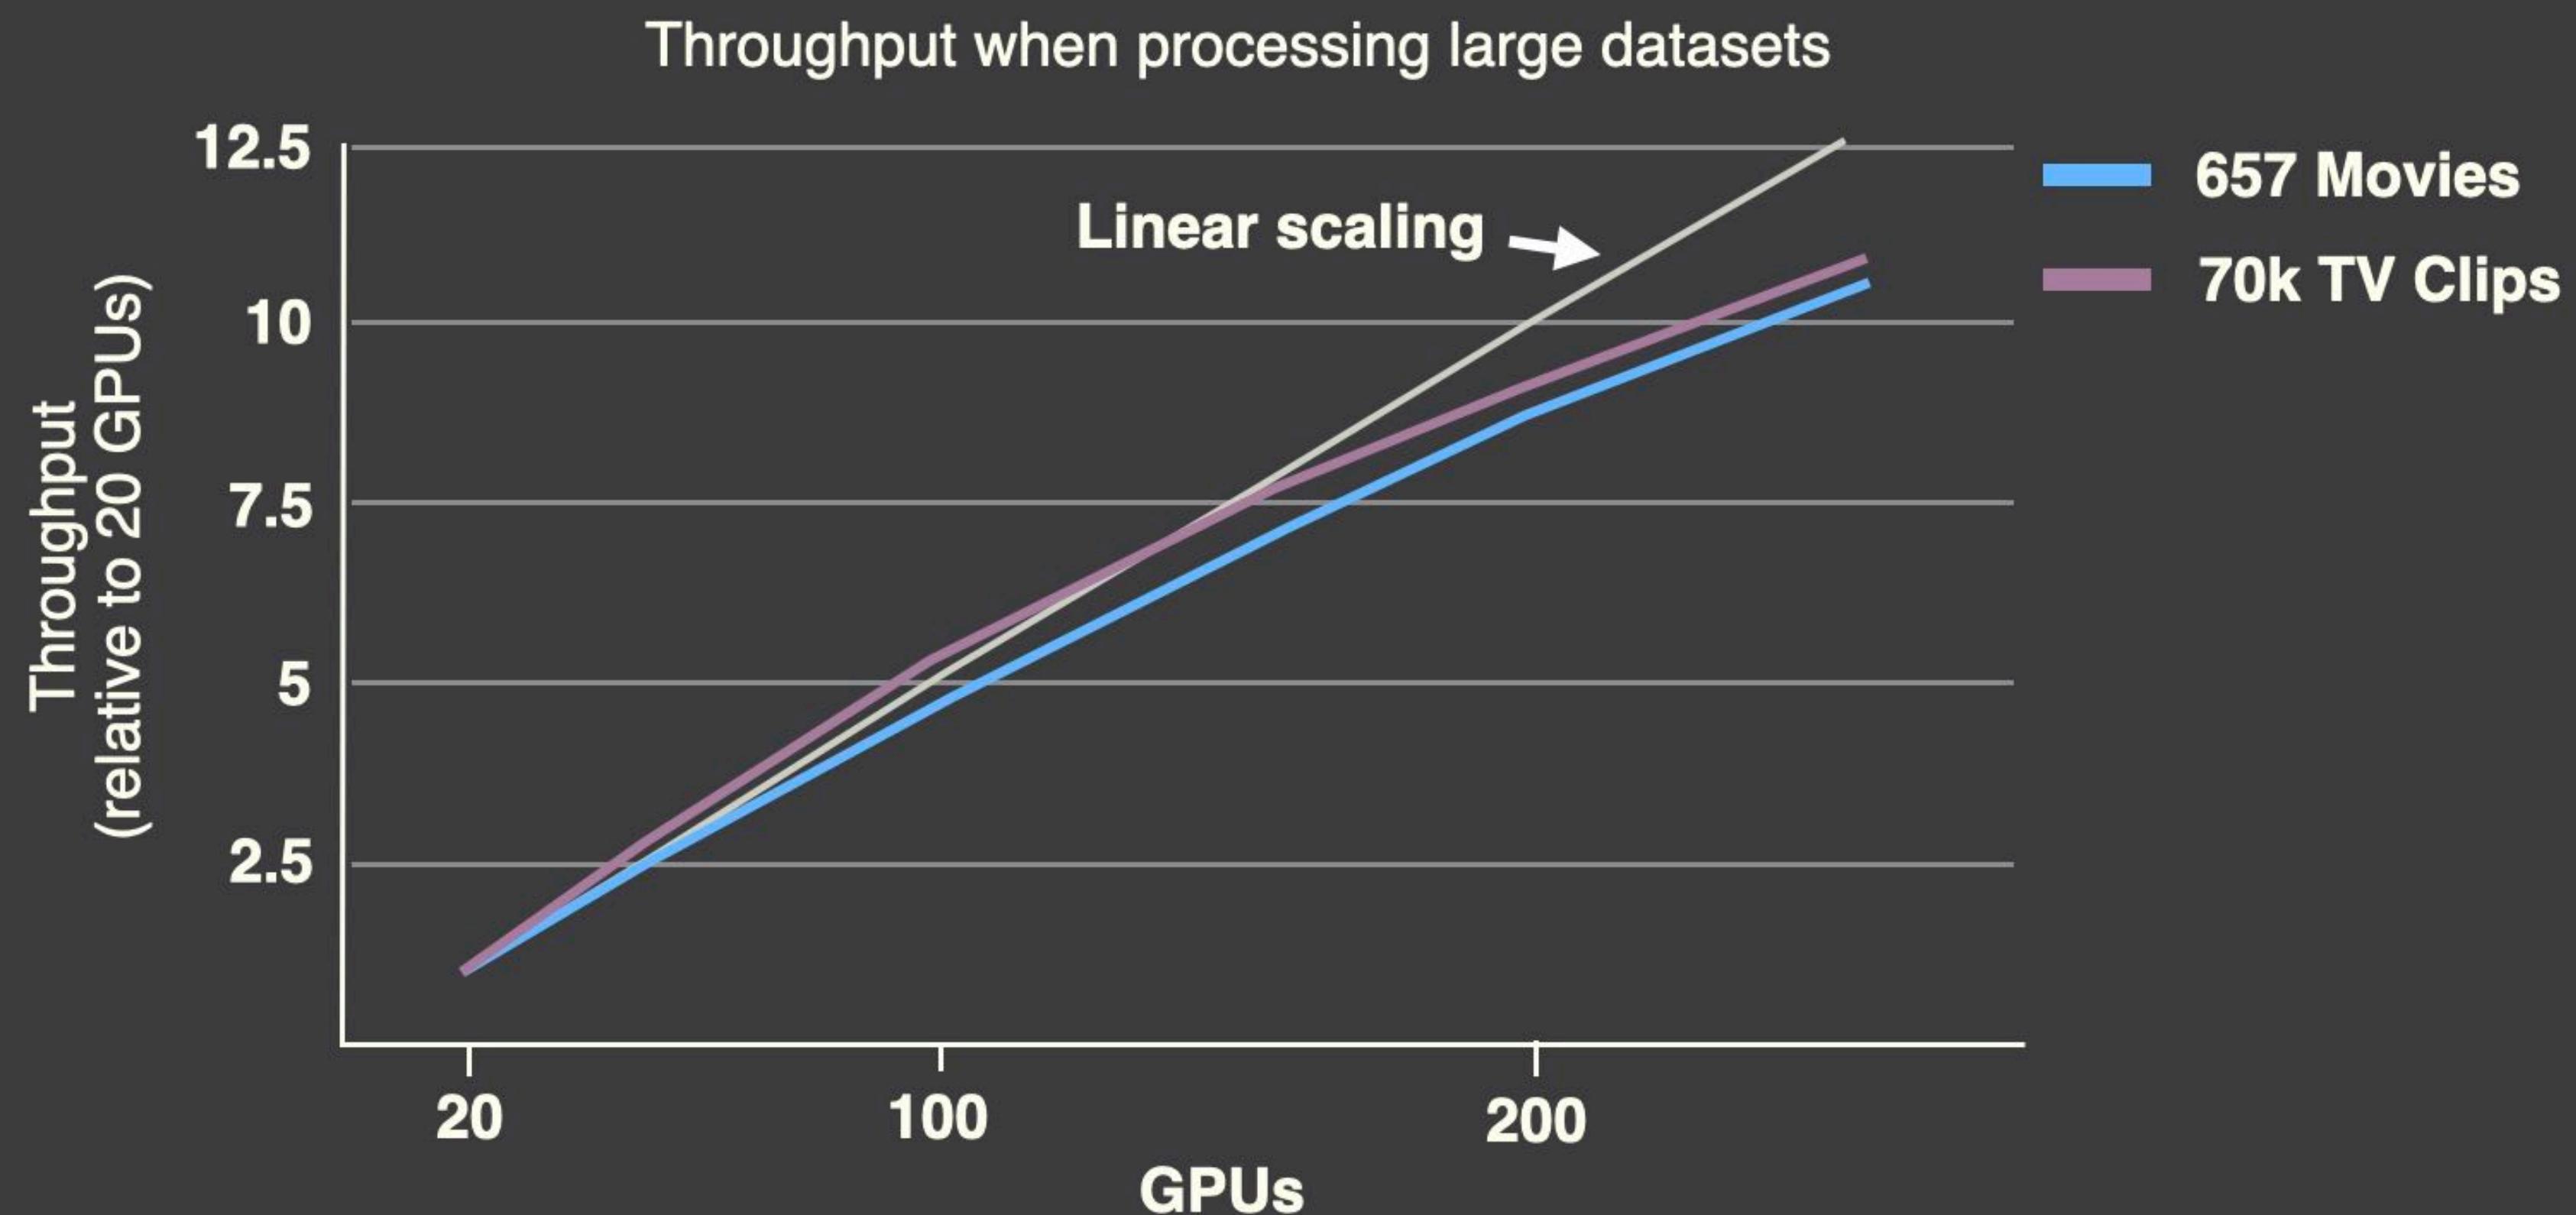
**One point per slide!**

**(and the point is the title of the slide!!!)**

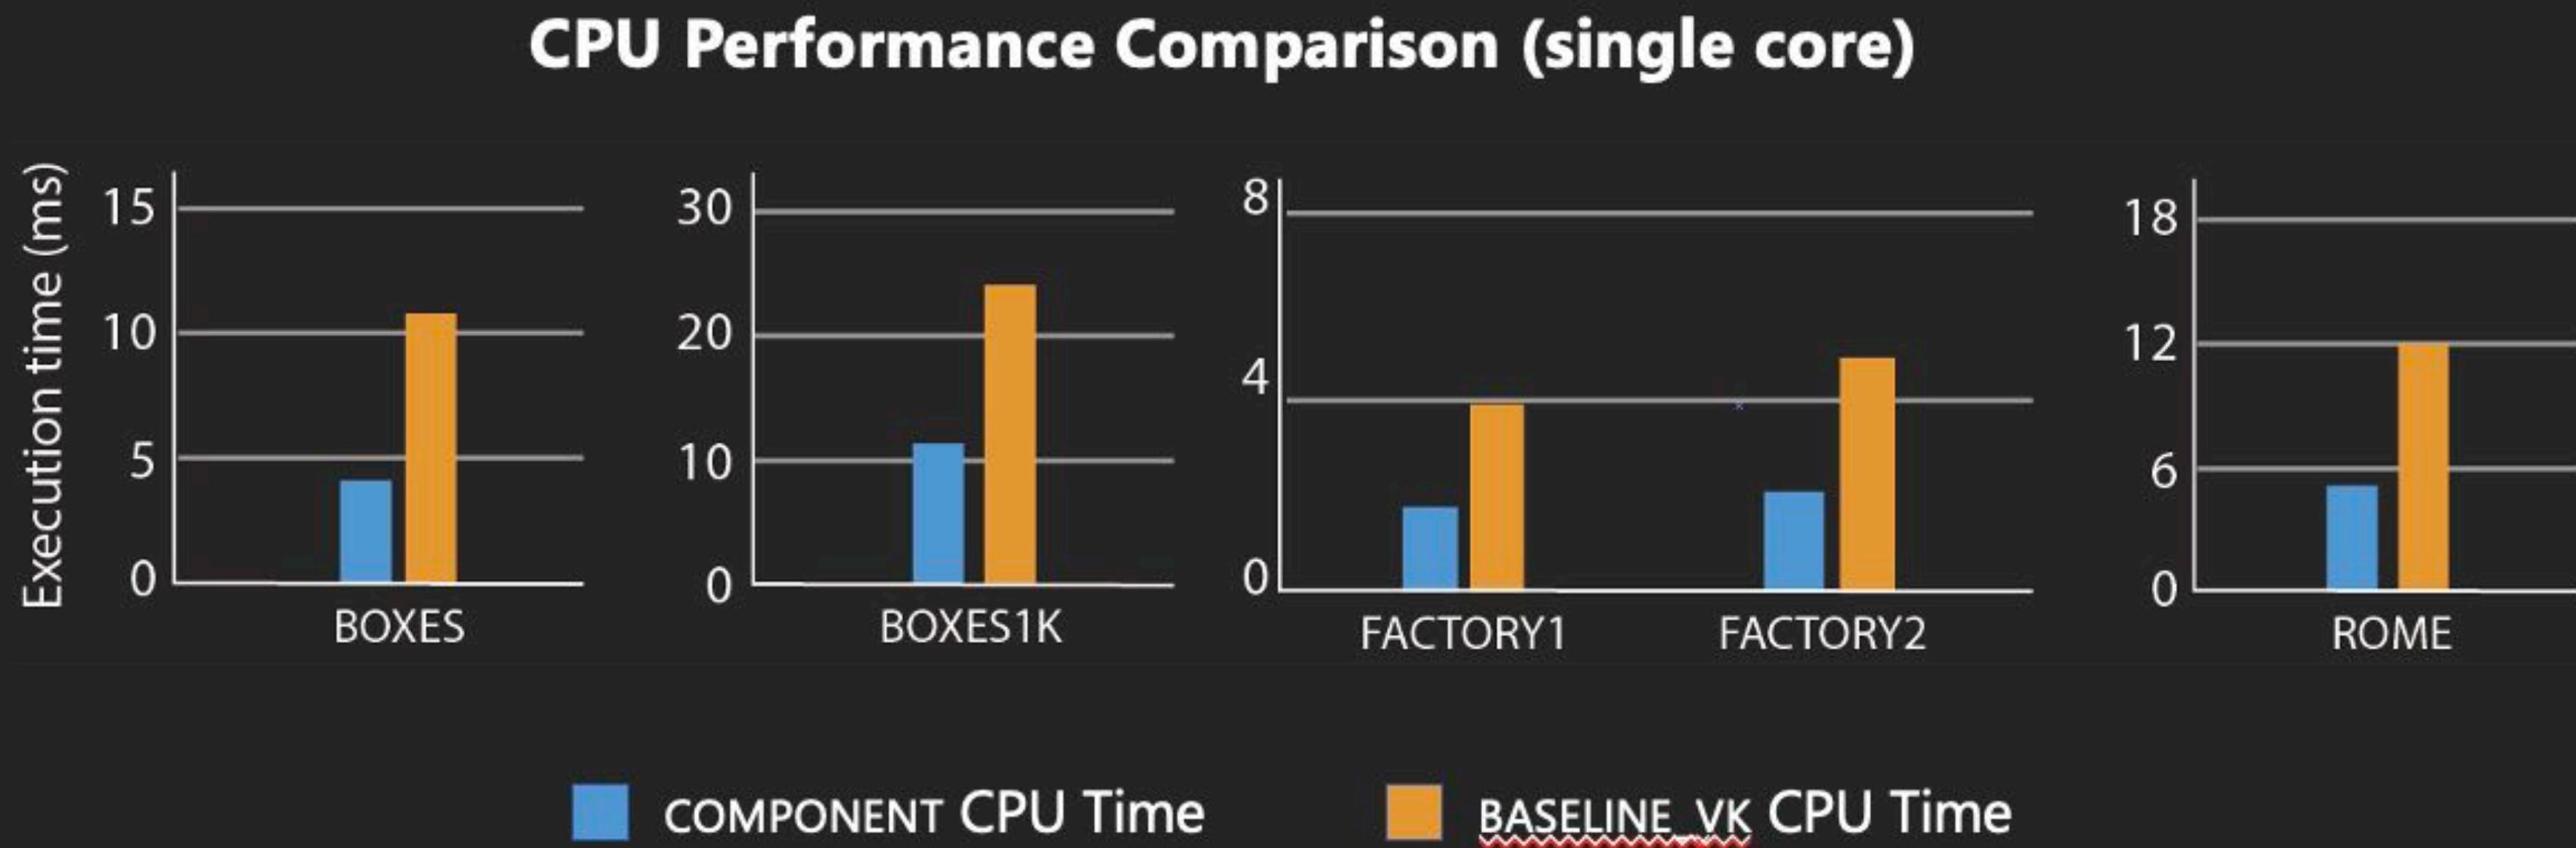
## ■ Make the point of the graph the slide's title:

- It provides audience context for interpreting the graph (“Let me see if I can verify that point in the graph to check my understanding”)
- Another example of the “audience prefers not to think” principle

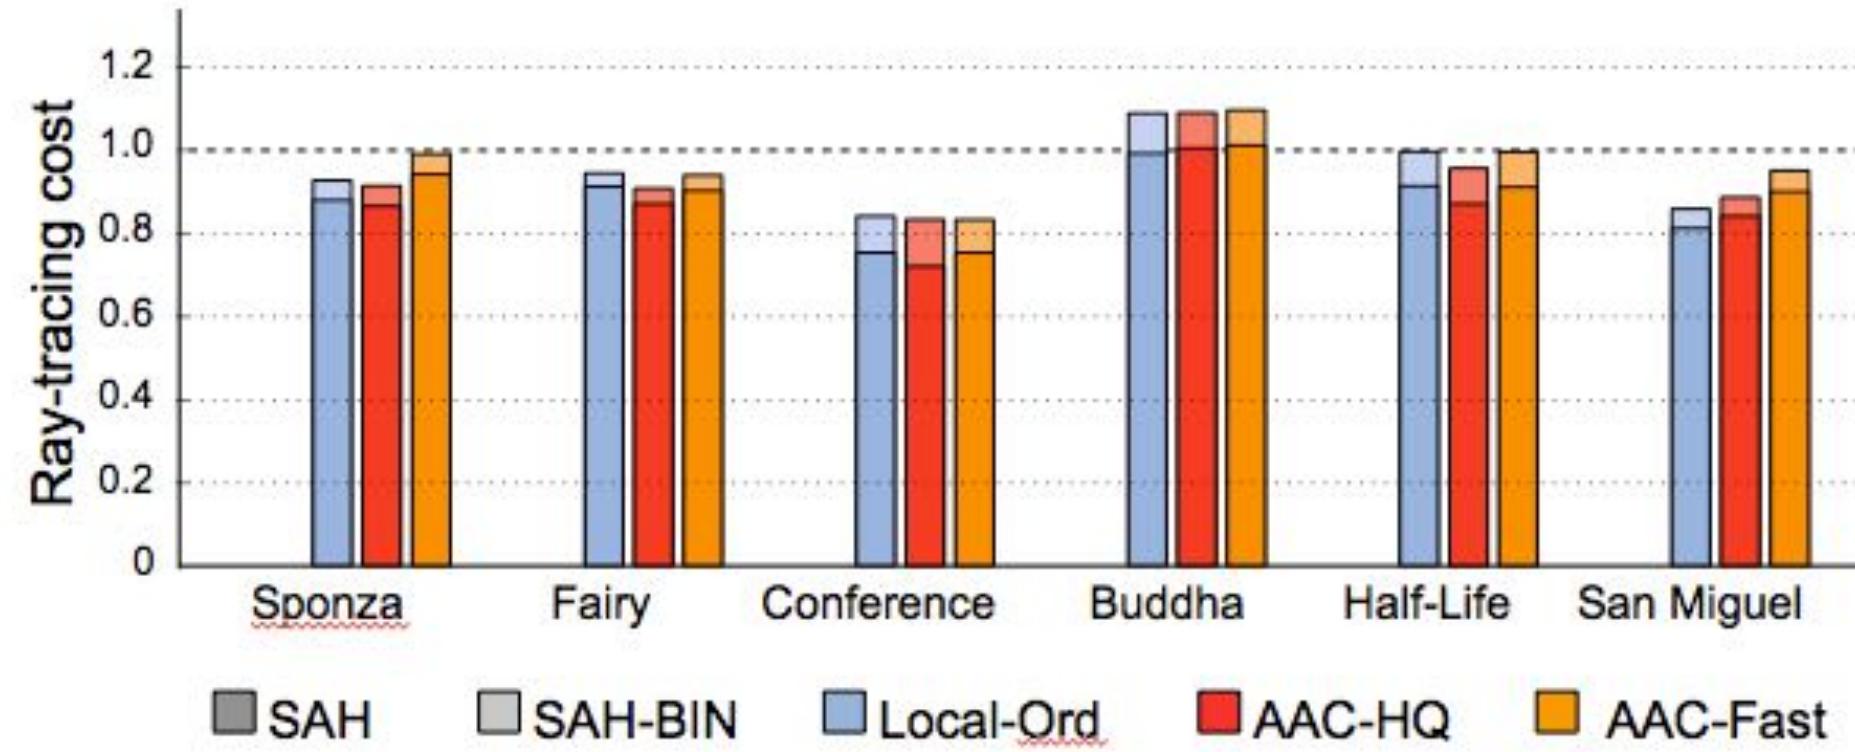
## Scanner scales when processing large datasets



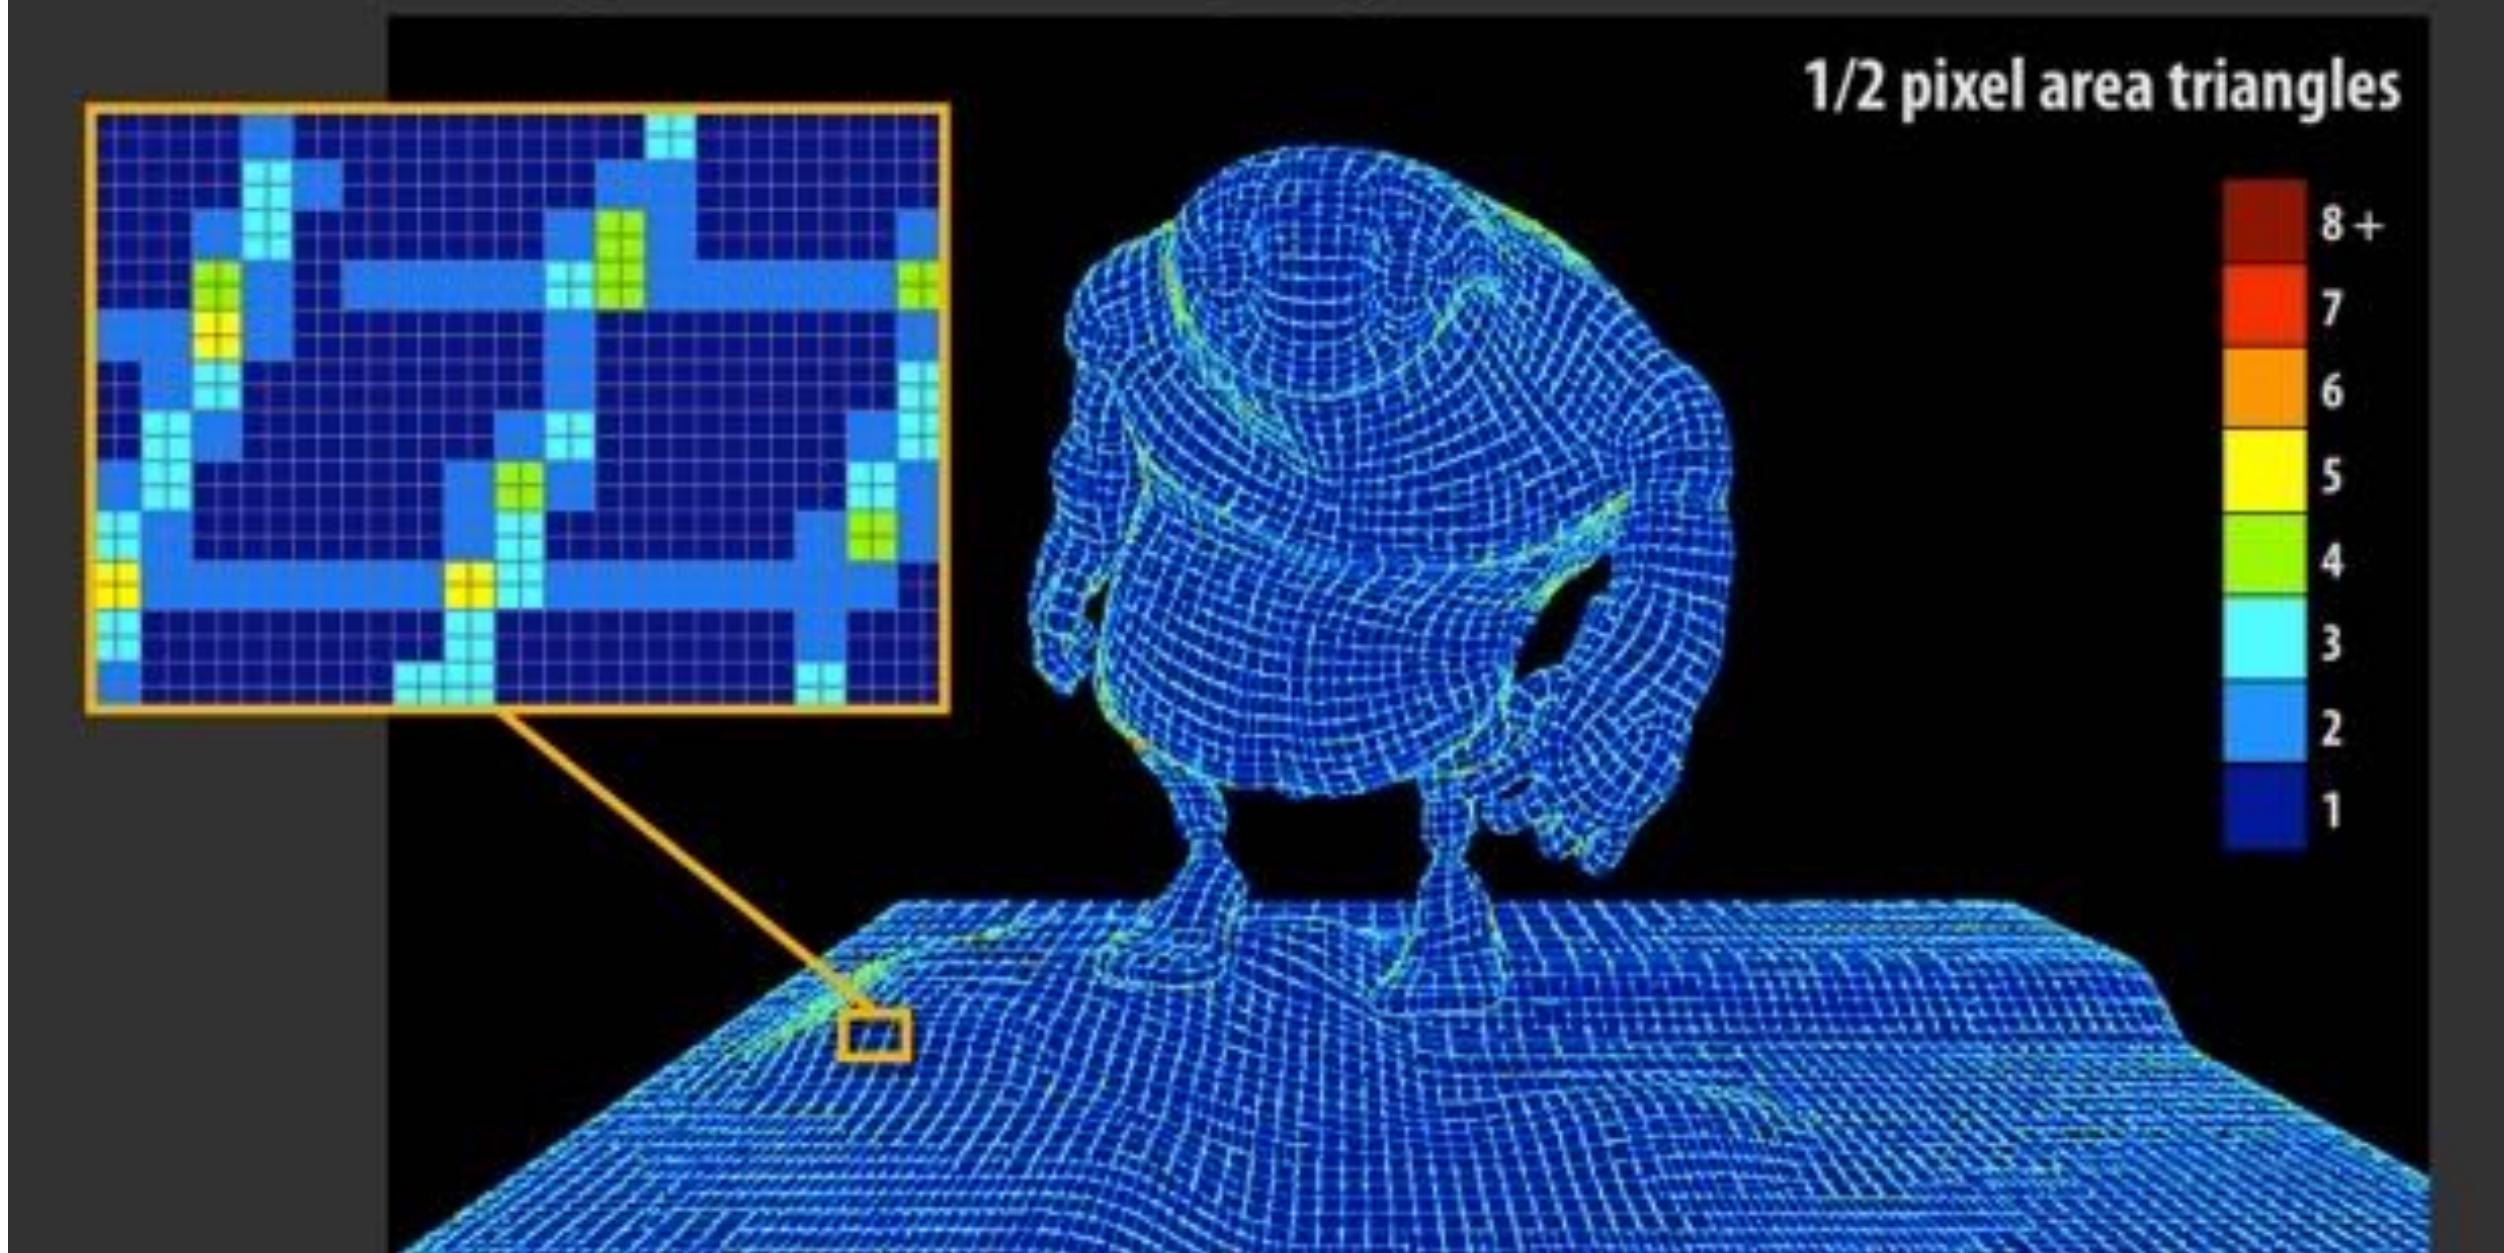
# The **COMPONENTS** renderer uses 2x less CPU time than **BASELINE\_VK**



AAC-Fast produces BVHs with equal or lower cost than the full sweep build in all cases except Buddha.



Extra shading occurs at merging window boundaries

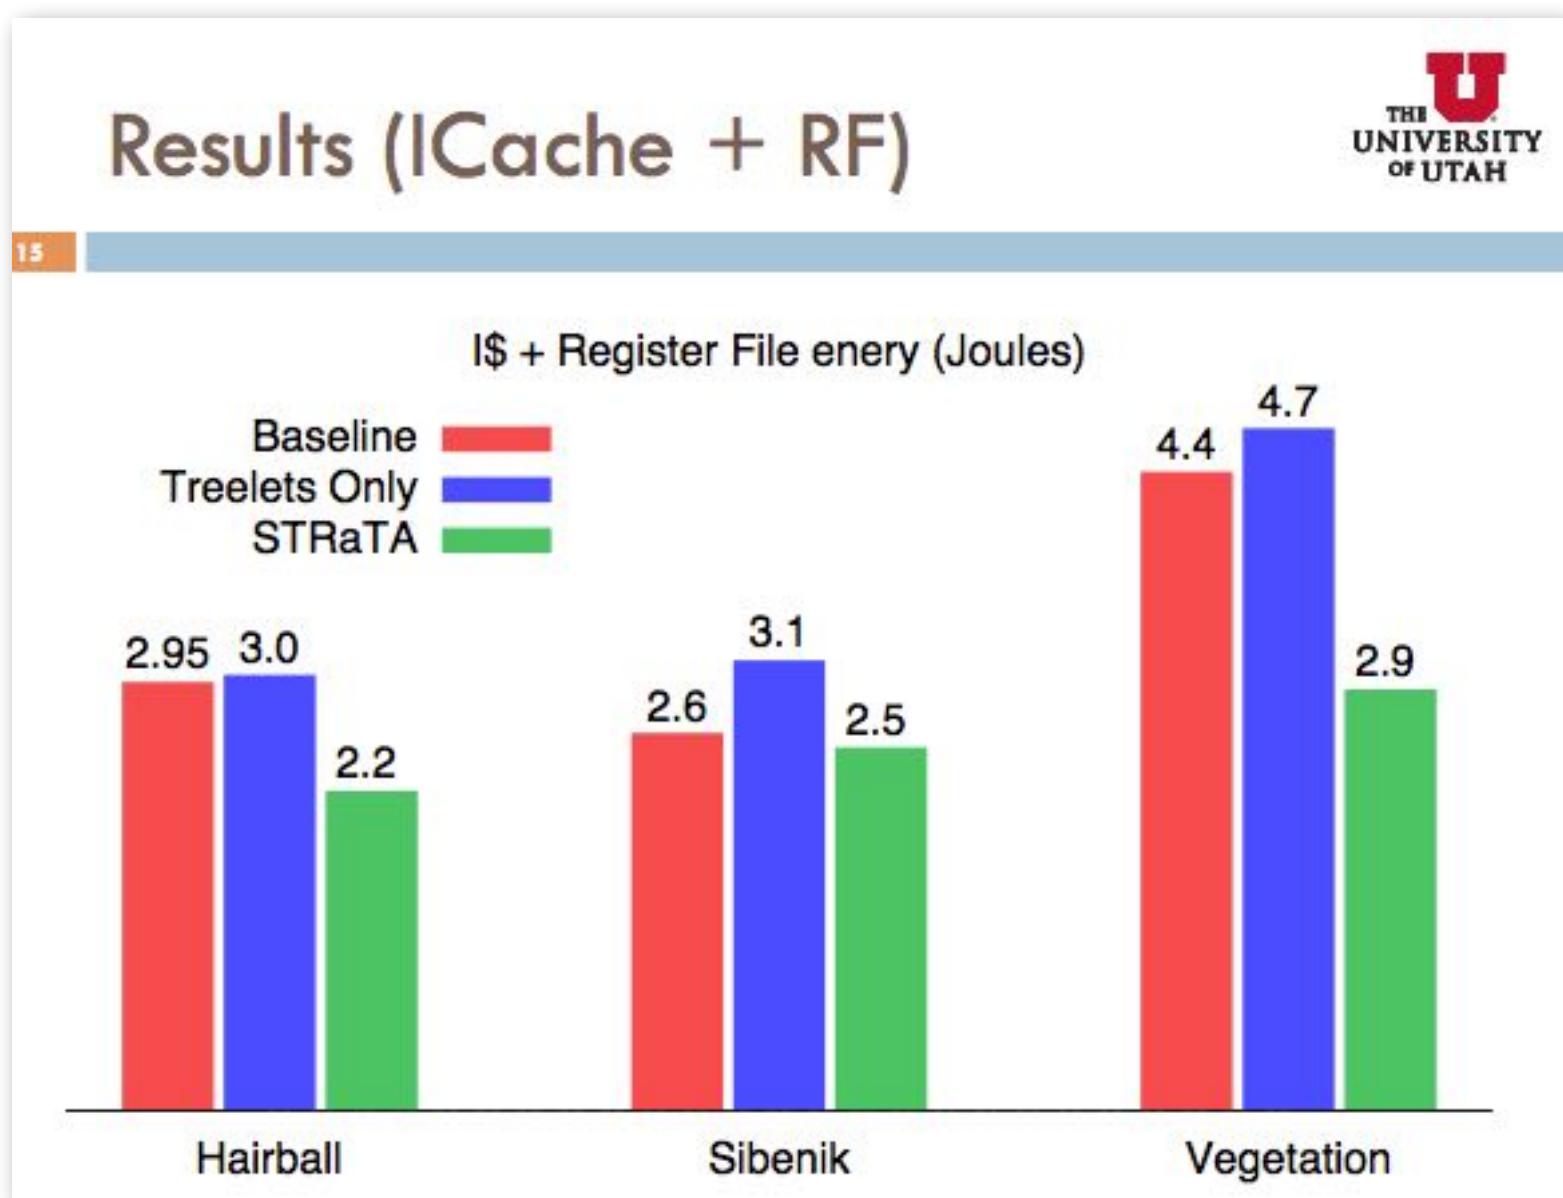


# Corollary to the one point per slide rule

- In general, you don't want to show data on a results slide that is unrelated to the point of the slide
- This usually means you need to remake the graphs from your paper (it's a pain, but sorry, it's important) \*
- This is the “every sentence matters” principle applied to visual details on a slide

\* This is an example of a tip for conference talk polish: not necessary for informal talks

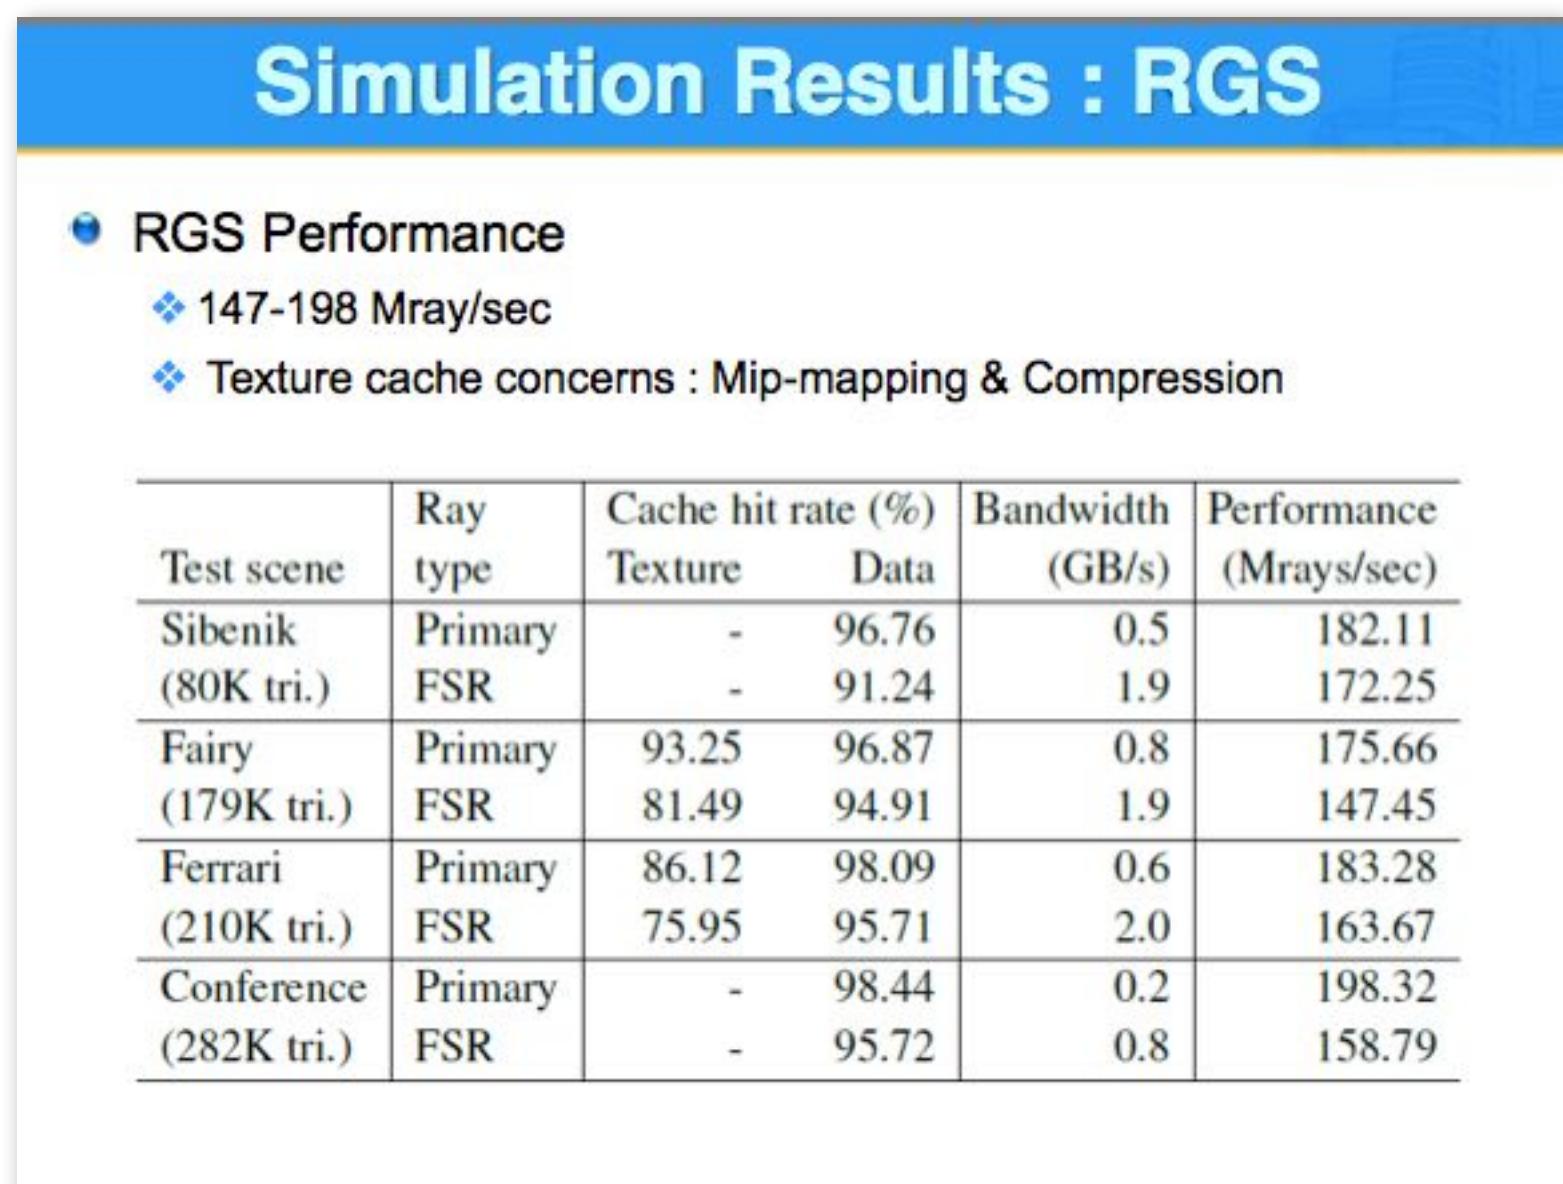
# Bad examples of results slides



- Notice how you (as an audience member) are working hard to interpret the trends in these graphs

- You are asking: what do these results say?

- You just want to be told what to look for



# Tip 9

## Titles matter

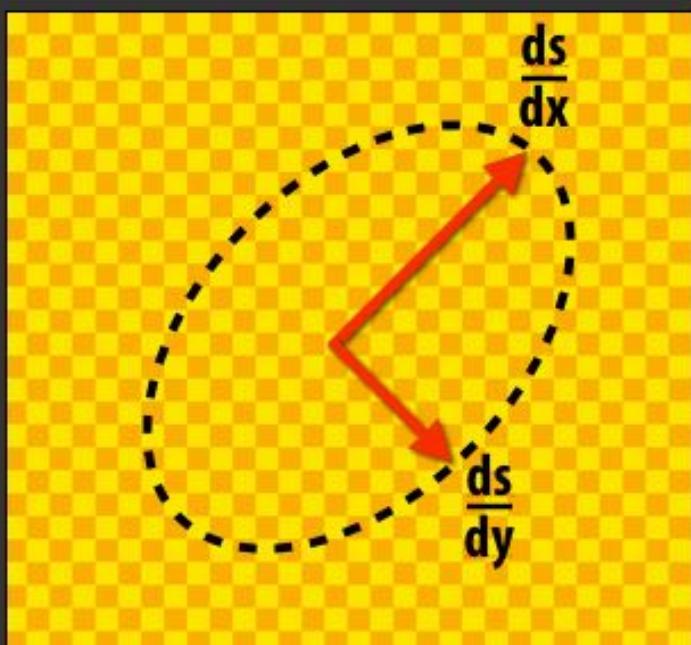
**If you read the titles of your talk all the way through, it should be a great summary of the talk.**

**(basically, this is “one-point-per-slide” for the whole talk)**

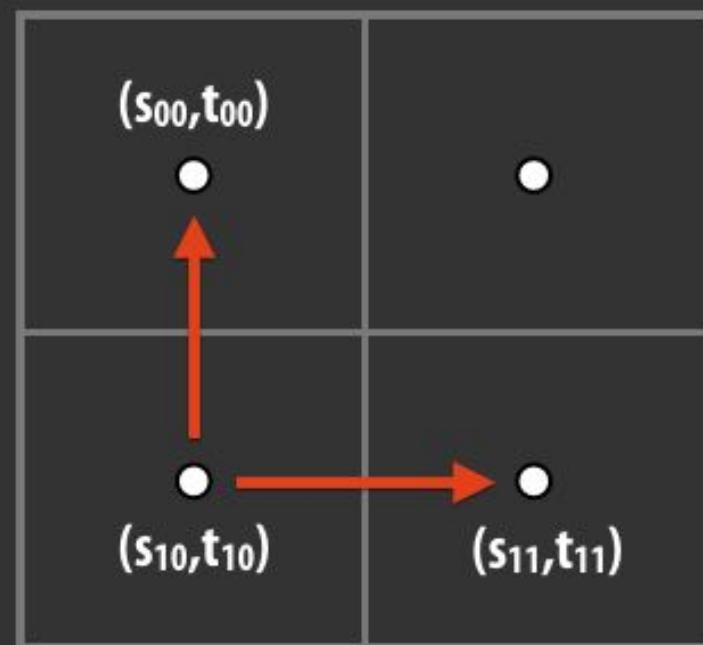
# Examples of good slide titles

## GPUs shade quad fragments (2x2 pixel blocks)

Texture data



Quad fragment



use differences between neighboring  
texture coordinates to estimate derivatives

## Greedy SRDH build optimizes over partitions and traversal policies

SAH:

```
forall(partitions in set-of-partitions)
    ...evaluate SAH and pick min...
```

SRDH:

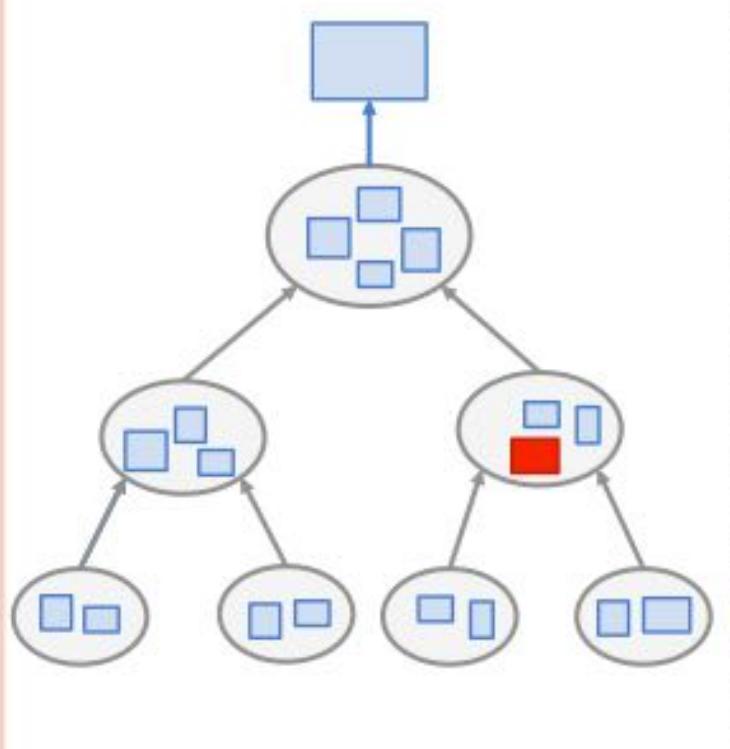
```
forall(partitions in set-of-partitions)
    forall(traversalKernels in set-of-kernels)
        ...evaluate SRDH and pick min...
```

$$\text{SRDH}(R,L,\kappa,r) = (1 - \kappa(r) H(L,r)) |R| + (1 - \kappa(r) H(R,r)) |L|$$

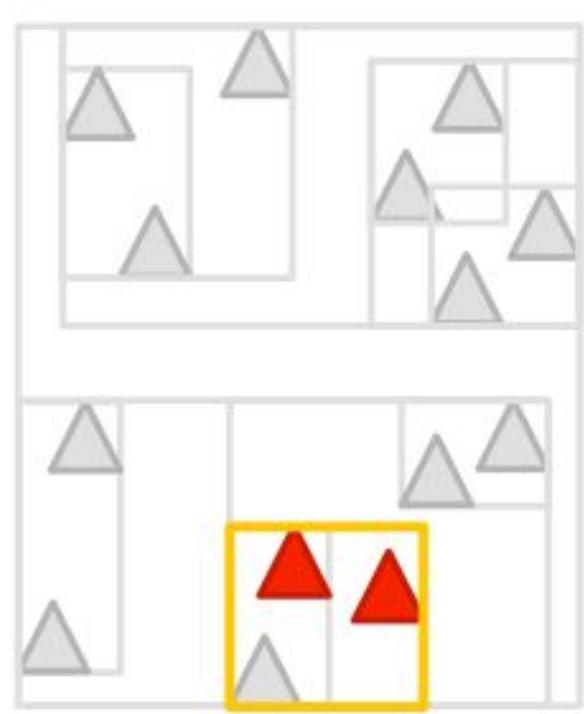
51

**AAC IS AN APPROXIMATION TO THE TRUE AGGLOMERATIVE CLUSTERING SOLUTION.**

Computation graph:



Primitive partitioning:



**The reason for meaningful slide titles is convenience and clarity for the audience**

**“Why is the speaker telling me this again?”**

**(Recall “why before what”)**

# Read your slide titles in thumbnail view

Do they make all the points of the story you are trying to tell?

The following is a summary of the content in each of the 36 slides shown in the grid:

- Slide 1:** Reducing Shading on GPUs using Quad-Fragment Merging. Shows authors: Kayvon Fatahalian, Soleiman Bealai, James Hinchey, Kurt Akeley, Pat Hanrahan, Michael R. Mark, and Henry Moreton.
- Slide 2:** High-resolution meshes are appearing in games. Shows a dragon and a house with "Low detail" and "High detail" versions.
- Slide 3:** High-resolution meshes are appearing in games. Shows a dragon and a house with red boxes highlighting edges.
- Slide 4:** PROBLEM: Current GPUs shade small triangles inefficiently.
- Slide 5:** Multi-sample locations. Shows a triangle with a grid of sample points and a color bar for "Sample coverage multiple times per pixel (for anti-aliased edges)".
- Slide 6:** Shading sample locations. Shows a triangle with a grid of sample points and a color bar for "Sample shading once per pixel".
- Slide 7:** Surface derivatives are needed for texture filtering. Shows a yellow textured surface with a red arrow indicating a derivative.
- Slide 8:** GPUs shade quad fragments (2x2 pixel blocks). Shows a yellow textured surface with a 2x2 grid of shaded pixels labeled "quad fragment".
- Slide 9:** Shaded quad fragments. Shows a yellow textured surface with a 2x2 grid of shaded pixels labeled "quad fragment".
- Slide 10:** Final pixel values. Shows a yellow textured surface with a 2x2 grid of final pixel values.
- Slide 11:** Pixels at triangle boundaries are shaded multiple times. Shows a yellow textured surface with a 2x2 grid of shaded pixels with a color bar for "Shading computations per pixel".
- Slide 12:** Pixels at triangle boundaries are shaded multiple times. Shows a yellow textured surface with a 2x2 grid of shaded pixels with a color bar for "Shading computations per pixel".
- Slide 13:** Pixels at triangle boundaries are shaded multiple times. Shows a yellow textured surface with a 2x2 grid of shaded pixels with a color bar for "Shading computations per pixel".
- Slide 14:** Small triangles result in extra shading. Shows three triangles with labels: "tiny overtriangle", "biggest triangle", and "1-pointed triangles".
- Slide 15:** Goal: Shade high-resolution meshes (not individual triangles) approximately once per pixel. Approach: Evolve GPU's quad-fragment shading system (Provide smooth evolution from status quo).
- Slide 16:** QUAD-FRAGMENT MERGING.
- Slide 17:** GPU pipeline (with tessellation). Shows the GPU pipeline stages: Vertex → Tessellation → Pixel → Shading.
- Slide 18:** Rasterized quad-fragment. Shows a 2x2 grid of shaded pixels.
- Slide 19:** Rasterized quad-fragment. Shows a 2x2 grid of shaded pixels with a "multi-sample coverage mask".
- Slide 20:** Rasterized quad fragments. Shows a 2x2 grid of shaded pixels with a "multi-sample coverage mask".
- Slide 21:** GPU pipeline: triangle connectivity is known. Shows the GPU pipeline stages with a callout to "Triangle connectivity is known".
- Slide 22:** Pipeline with quad-fragment merging. Shows the GPU pipeline stages with a callout to "Pipeline with quad-fragment merging".
- Slide 23:** Pipeline with quad-fragment merging. Shows the GPU pipeline stages with a callout to "Pipeline with quad-fragment merging".
- Slide 24:** Two key merging operations:
  - Identifying when quad fragments can be merged
  - Constructing a merged quad fragment
- Slide 25:** Merging quad fragments. Shows a mesh triangle, rasterized quad fragments, and merged quad fragments.
- Slide 26:** Merging quad fragments. Shows a mesh triangle, rasterized quad fragments, and merged quad fragments.
- Slide 27:** Merging quad fragments. Shows a mesh triangle, rasterized quad fragments, and merged quad fragments.
- Slide 28:** Two key merging operations:
  - Identifying when quad fragments can be merged
  - Constructing a merged quad fragment
- Slide 29:** Challenge: Avoiding merges that introduce visual artifacts.
- Slide 30:** Example: surface with a silhouette. Shows a surface with a silhouette and a resulting shaded result with anti-aliased silhouettes.
- Slide 31:** Naive merging results in aliasing. Shows a mesh triangle, rasterized quad fragments, and a "aliased result".
- Slide 32:** Avoid merging across discontinuities. Shows a mesh triangle, rasterized quad fragments, and a "mesh tiling" diagram.
- Slide 33:** Conditions required to merge quad fragments:
  - Same screen location
  - Same sidedness (triangles front-facing or back-facing)
  - Source triangles are adjacent in the mesh
- Slide 34:** High-frequency geometric detail may cause aliasing. Shows a wavy line with a callout to "Our merging rules are designed for real-time performance".
- Slide 35:** Implementation: the cost of merging is low:
  - Merging operations are cheap
    - Testing merging rules requires only integer operations
  - Merge buffer is small
    - 12 quad-fragment merge buffer is very efficient
  - Expectation: quad-fragment merging can be encapsulated in fixed-function hardware
- Slide 36:** EVALUATION.

# **Tip 10**

**Organize your talk with section slides**

I'm about to frame the problem in my terms

**Reducing Shading on GPUs using Quad-Fragment Merging**

**High-resolution meshes are appearing in games**

**PROBLEM**  
Current GPUs shade small triangles inefficiently

**Multi-sample locations**

**Shading sample locations**

**Surface derivatives are needed for texture filtering**

**GPUs shade quad fragments (2x2 pixel blocks)**

**Shaded quad fragments**

**Final pixel values**

**Pixels at triangle boundaries are shaded multiple times**

**Pixels at triangle boundaries are shaded multiple times**

**Small triangles result in extra shading**

**Goal:**  
Shade high-resolution meshes (not individual triangles) approximately once per pixel

**Approach:**  
Evolve GPU's quad-fragment shading system (Provide smooth evolution from status quo)

**GPU pipeline (with tessellation)**

**Rasterized quad-fragment**

**Rasterized quad-fragment**

**Rasterized quad fragments**

**GPU pipeline: triangle connectivity is known**

**Pipeline with quad-fragment merging**

**Pipeline with quad-fragment merging**

**Two key merging operations**

**Merging quad fragments**

**Merging quad fragments**

**Merging quad fragments**

**Challenge**  
Avoiding merges that introduce visual artifacts

**Example: surface with a silhouette**

**EVALUATION**

**Naive merging results in aliasing**

**Avoid merging across discontinuities**

**Conditions required to merge quad fragments**

**High-frequency geometric detail may cause aliasing**

**Implementation: the cost of merging is low**

**Experimental setup**

**Merging reduces total shaded quad fragments**

**Merging reduces total shaded quad fragments**

**Extra shading occurs at merging window boundaries**

**Form micropolygons: factor of eight across scenes**

**Rough surfaces result in less merging**

**Nearly identical visual quality**

**Differences exist near silhouettes**

**Shader derivatives alias in areas of high curvature**

**High visual quality for smooth surfaces**

**ALTERNATIVES**

**Reyes**

**Quad-fragment merging**

**A real-time micropolygon rendering pipeline ...**

**Let me wrap up**

# Stage your talk with section slides

## ■ Useful for your audience

- It provides guidance for what you hope to achieve next (no surprises!)
- Compartmentalization: it's absolutely clear where the shifts are
- If a listener got lost, it's a good place for them to re-engage
- It's a place for the audience to take a breath
- It gives the talk a more colloquial tone

## ■ Useful for you

- It's a chance for you to pause and take a breath
- It's a great breakdown of the talk for practicing subsections

# **Tip 11**

**End on a positive note!**

# End on a positive note!

**The future is bright!**

**Lots of new work to do, here are some ideas!**

**This is one part of something bigger!**

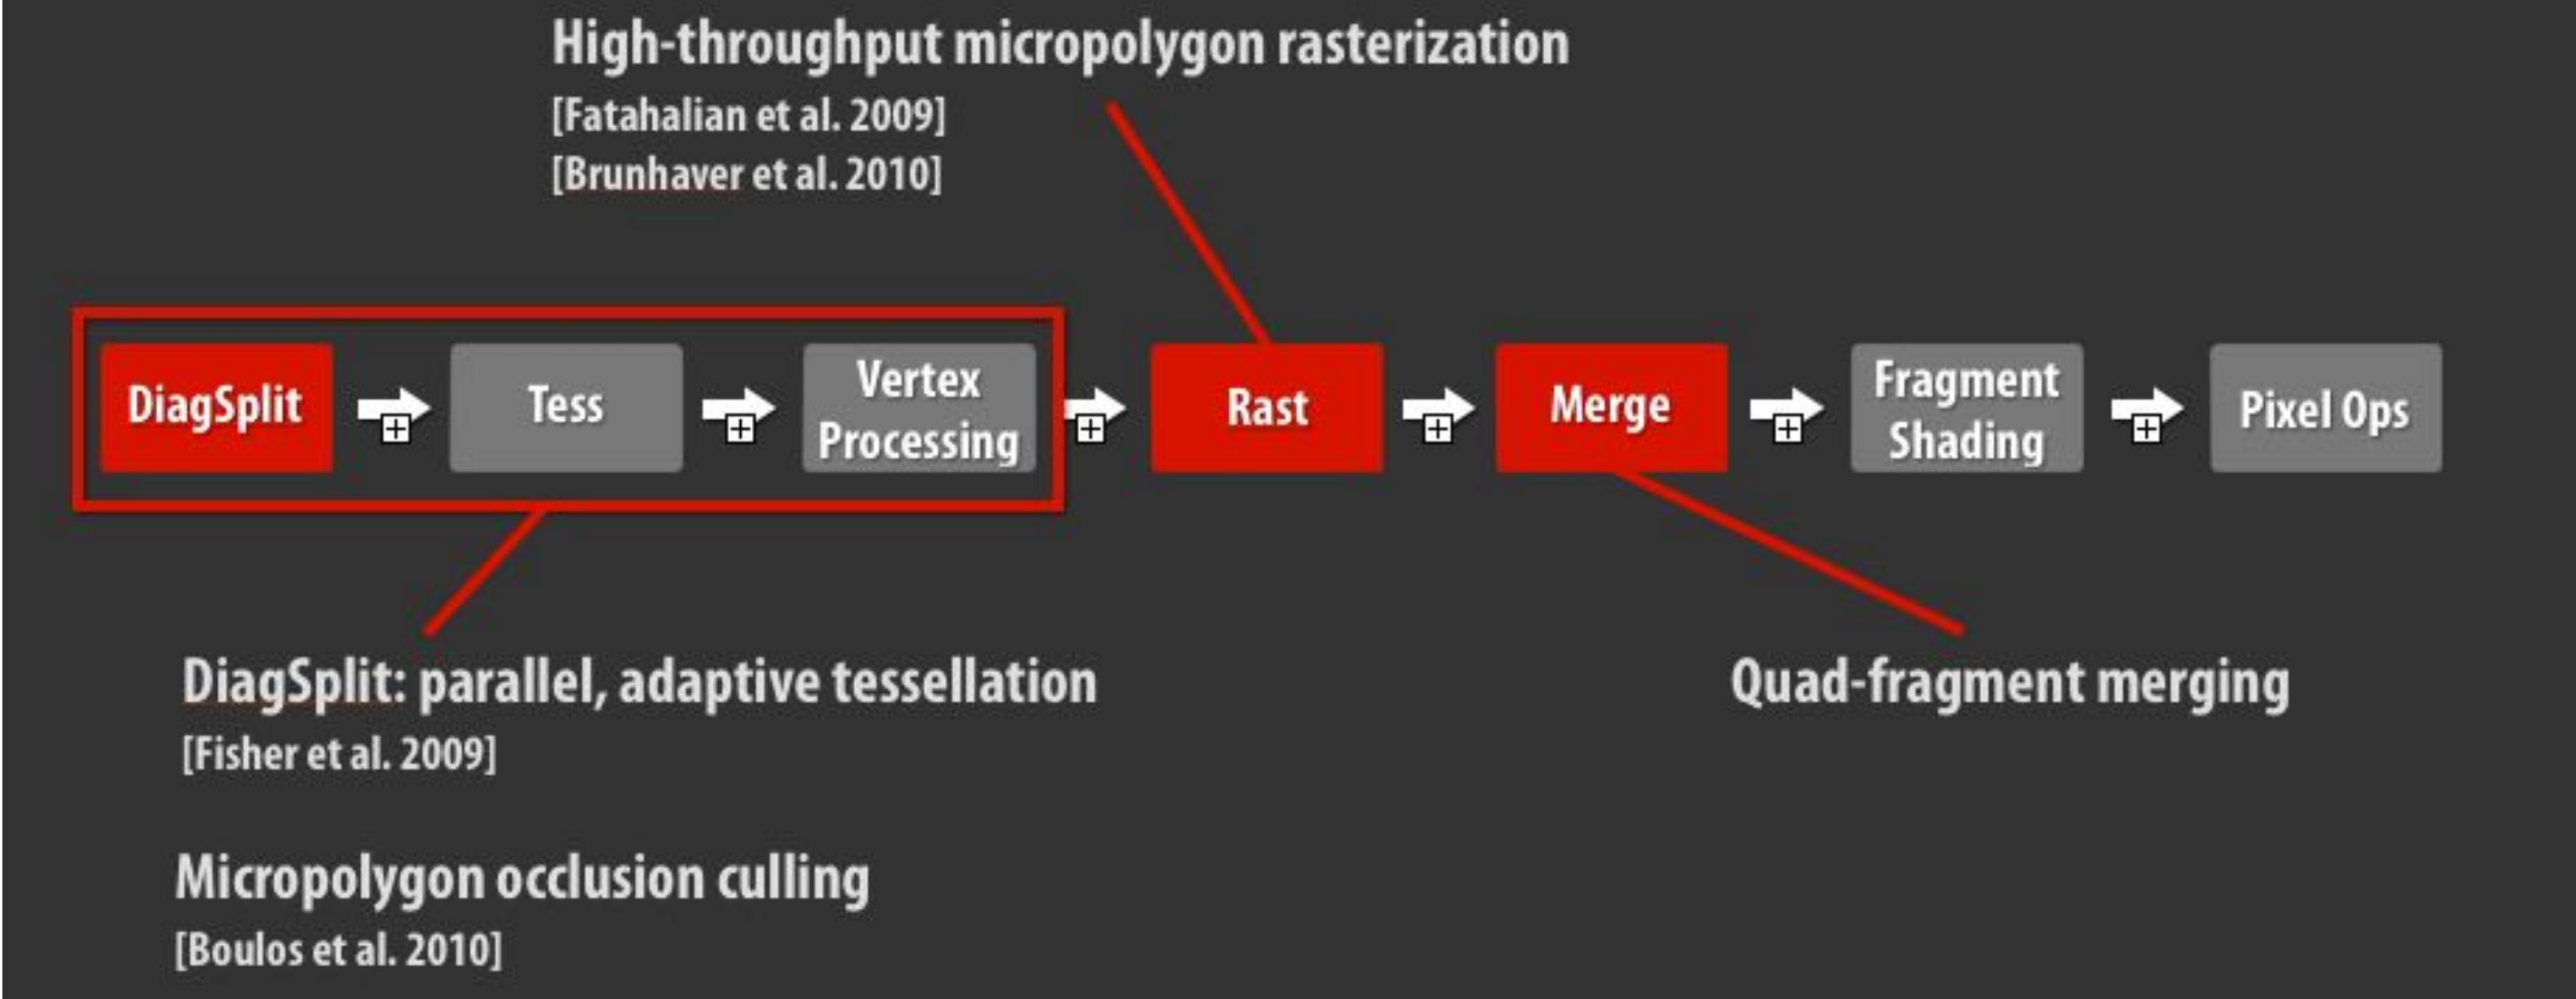
**Many talks end on future work in a manner that stresses problems with the current work or enumerates obvious next steps**

- **It's boring and sort of a bummer for everyone involved.**  
**(Audience: "well, that was a bit incremental.")**
- **It's a lost opportunity to impart critical intellectual thought to the field \***
  - **Recall: introduction was where you contributed critical intellectual thought in how to think about the problem being solved today.**
  - **Conclusion is where you contribute intellectual thought that reflects on what you have done, or about the future.**

\* Every sentence matters

# Final SIGGRAPH slide as a Ph.D. student

A real-time micropolygon rendering pipeline ...  
is not far away.



An earlier draft had a very simple future work slide (“we could do X, Y, Z”). I was told by my advisor that it was a let down and to think about how to end on a broader note. This slide took about four hours to come up with.

Result: Tony DeRose got it! He realized the point wasn’t just the one particular optimization that was the contribution of the SIGGRAPH paper, but a broader line of work on rethinking the graphics pipeline for high-quality rendering: his comment was of the effect, “I’m glad someone’s finally figured the big pieces of this out.”

# **Tip 12**

**The audience is always right:  
When receiving feedback on a practice talk,  
do not be defensive!**

# The audience is always right

- Your tendency will be to be defensive when someone claims an idea in your talk was not explained well or was not clear
- You will find yourself turning to the relevant slide in your talk and saying “I mentioned that here”.
- The customer (the audience) is right in this situation. Sure, you might have mentioned it, but if it wasn’t understood, it’s your fault not theirs.
  - Find a way to make it more clear!
- The correct response is to turn to the appropriate slide and say:
  - “I tried to explain that idea here. And this is what I was thinking. What could I have said to make that point more clear?”
- The complainer should then work with you to explain what they interpreted instead, and offer suggestions on what information they would require to have better understood.

# **Wrap Up**

# **General principles to keep in mind**

**Identify your audience, and strive for perfect clarity for them.**

**“Every sentence matters.”**

**“Show, don’t tell.”**

**“The audience prefers not to think” (about things you can just tell them)**

**“Surprises are bad”: say why before what**

**(indicate why you are saying something before you say it)**

**Explain every figure, graph, or equation**

**When improving the talk, the audience is always right**

# Useful reminders

## ■ In general:

- Every sentence matters. Pick your target audience level and aim for them to understand everything. If it's not clear to this audience, take it out.

## ■ In your introduction

- Your goal is to tell the audience how you want them to think about the problem your talk is about.

## ■ In the remainder of the talk (in particular, the evaluation)

- Always explain every figure, graph, or equation
- One point per slide
- Place the point in title of slide

## **Summary...**

**Some very simple principles enable much more efficient, and much more rigorous, communication about technical content.**



**Take pride in your craft. All students should follow these principles until they become second nature, and then you can start violating them.**

